

Robustness under pressure.

A measured assessment of an emotion classifier before and after application-layer hardening.

V1 / Unhardened

V2 / Four defenses

Analysis schema 2

Model: j-hartmann/emotion-english-distilroberta-base · 7 emotion labels

THE COMPARISON

What changed after hardening?

Compatible inputs

Recorded coverage · V1 1928/1928 (100.00%) · V2 1928/1928 (100.00%)

15

Resolved groups

Original evidence no longer shows that failure mode.

14

Remaining groups

Positive evidence of the same weakness in V2.

0

Inconclusive V1 groups

Insufficient evidence to establish resolution.

0

Newly demonstrated groups

Usable V1 evidence establishes a V2 regression.

Groups are category × subfamily × failure mode. Counts describe weaknesses, not individual requests or overall security.

Where the pressure landed.

Recorded coverage · V1 1928/1928 (100.00%) · V2 1928/1928 (100.00%)

Failure rate by attack category

● V1 ● V2

Failed / eligible automatically scored attacks. Each attack counts once per category, even if it triggers several failure modes. Rates reflect this suite's composition.

CATEGORY	V1 · UNHARDENED		V2 · HARDENED	
Malformed	19.05%	4 / 21	4.76%	1 / 21
Boundary	9.09%	1 / 11	0.00%	0 / 11
Perturbation	11.94%	45 / 377	11.94%	45 / 377
Encoding	11.17%	104 / 931	6.77%	63 / 931
Whitespace	32.38%	68 / 210	0.00%	0 / 210
Truncation	52.38%	88 / 168	2.38%	4 / 168

Operational observations

Scored attacks only. A request error, timeout, and failed health check are separate observations and can overlap.

OUTCOME	V1	V2
Unhandled HTTP 5xx	91	0
Request timeouts	0	0
Connection failures	0	0
Failed health checks	0	0
Slow responses	1	1
Information leaks	0	0
Invalid input accepted	1	0

V1 Unhardened endpoint

RECORDED COVERAGE

1928 / 1928

100.00% of planned cases recorded

Selected	1928
Missing selected cases	0
Skipped	0
Excluded by limit	0

QUALIFYING PREDICTION FLIPS

15.91%

218 / 1370 eligible prediction pairs

36 of 42 baselines eligible for flip scoring

Top label changed and clean confidence ≥ 0.6 . Each version uses its own clean prediction.

Mean confidence change on qualifying flips: -0.2414 (probability units; not accuracy).

[View V1 denominators, exclusions and per-flip evidence in the appendix ↓](#)

V2 Hardened endpoint

RECORDED COVERAGE

1928 / 1928

100.00% of planned cases recorded

Selected	1928
Missing selected cases	0
Skipped	0
Excluded by limit	0

QUALIFYING PREDICTION FLIPS

8.16%

112 / 1373 eligible prediction pairs

36 of 42 baselines eligible for flip scoring

Top label changed and clean confidence ≥ 0.6 . Each version uses its own clean prediction.

Mean confidence change on qualifying flips: -0.1439 (probability units; not accuracy).

[View V2 denominators, exclusions and per-flip evidence in the appendix ↓](#)

+ Unmatched case coverage

1928 matched recorded cases. 0 V1-only; 0 V2-only; 0 unavailable.

The evidence behind each weakness.

Recorded coverage · V1 1928/1928 (100.00%) · V2 1928/1928 (100.00%)

Findings are ordered by the supplied severity tier and score. Severity uses the strongest supporting instance; affected-case counts describe reproducibility. A resolved group concerns one failure mode only.

V1 finding register 29 groups

F-v1-009 · malformed/oversized_100kb · unhandled 5xx	F-v1-020 · perturbation/word_split · prediction flip
F-v1-012 · malformed/oversized_1mb · unhandled 5xx	F-v1-028 · whitespace/tabs · prediction flip
F-v1-001 · boundary/length_plus_one · unhandled 5xx	F-v1-025 · whitespace/interior_runs · prediction flip
F-v1-003 · encoding/compatibility_fullwidth · unhandled 5xx	F-v1-017 · perturbation/keyboard_type_second_word · prediction flip
F-v1-011 · malformed/oversized_10mb · unhandled 5xx	F-v1-004 · encoding/deletion_backspace · prediction flip
F-v1-021 · truncation/signal_end_over_limit · unhandled 5xx	F-v1-007 · encoding/invisible_zero_width · prediction flip
F-v1-023 · truncation/signal_start_over_limit · unhandled 5xx	F-v1-018 · perturbation/misspell · prediction flip
F-v1-002 · encoding/compatibility_fullwidth · prediction flip	F-v1-027 · whitespace/newlines · prediction flip
F-v1-005 · encoding/homoglyph · prediction flip	F-v1-029 · whitespace/unicode_spaces · prediction flip
F-v1-006 · encoding/homoglyph_greek · prediction flip	F-v1-026 · whitespace/leading_trailing · prediction flip
F-v1-013 · perturbation/capitalization · prediction flip	F-v1-022 · truncation/signal_end_short · prediction flip
F-v1-014 · perturbation/delete_char · prediction flip	F-v1-024 · truncation/signal_start_short · prediction flip
F-v1-015 · perturbation/insert_char · prediction flip	F-v1-008 · malformed/extra_field · invalid input accepted
F-v1-016 · perturbation/keyboard_type · prediction flip	F-v1-010 · malformed/oversized_10mb · slow response
F-v1-019 · perturbation/swap_adjacent · prediction flip	

F-v1-009 **Critical · 90/100** **Resolved**

malformed/oversized_100kb: unhandled server error

Malformed / oversized_100kb · unhandled 5xx · 1 affected case · GOLD

1 case(s) in malformed/oversized_100kb produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 1 case

malformed.oversized_100kb;

V1 → V2 evidence · 0 remaining / 1 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 1 case

malformed.oversized_100kb;

F-v1-012 **High · 80/100** **Resolved**

malformed/oversized_1mb: unhandled server error

Malformed / oversized_1mb · unhandled 5xx · 1 affected case · GOLD

1 case(s) in malformed/oversized_1mb produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 1 case

malformed.oversized_1mb;

V1 → V2 evidence · 0 remaining / 1 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 1 case

malformed.oversized_1mb;

F-v1-001 **High · 70/100** **Resolved**

boundary/length_plus_one: unhandled server error

Boundary / length_plus_one · unhandled 5xx · 1 affected case · GOLD

1 case(s) in boundary/length_plus_one produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 1 case

boundary.length_plus_one;

V1 → V2 evidence · 0 remaining / 1 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 1 case

boundary.length_plus_one;

F-v1-003 **High · 70/100** **Resolved**

encoding/compatibility_fullwidth: unhandled server error

Encoding / compatibility_fullwidth · unhandled 5xx · 3 affected cases · SILVER

3 case(s) in encoding/compatibility_fullwidth produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 3 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 3 cases

dataset-anger-5.encoding.compatibility_fullwidth.whole; dataset-joy-1.encoding.compatibility_fullwidth.whole; dataset-surprise-3.encoding.compatibility_fullwidth.whole;

V1 → V2 evidence · 0 remaining / 3 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 3 cases

dataset-anger-5.encoding.compatibility_fullwidth.whole; dataset-joy-1.encoding.compatibility_fullwidth.whole; dataset-surprise-3.encoding.compatibility_fullwidth.whole;

F-v1-011 **High · 70/100** **Resolved**

malformed/oversized_10mb: unhandled server error

Malformed / oversized_10mb · unhandled 5xx · 1 affected case · GOLD

1 case(s) in malformed/oversized_10mb produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 1 case

malformed.oversized_10mb;

V1 → V2 evidence · 0 remaining / 1 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 1 case

malformed.oversized_10mb;

truncation/signal_end_over_limit: unhandled server error

Truncation / signal_end_over_limit · unhandled 5xx · 42 affected cases · GOLD

42 case(s) in truncation/signal_end_over_limit produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 42 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 42 cases

dataset-anger-0.truncation.signal_end_over_limit; dataset-anger-1.truncation.signal_end_over_limit; dataset-anger-2.truncation.signal_end_over_limit; dataset-anger-3.truncation.signal_end_over_limit; dataset-anger-4.truncation.signal_end_over_limit; dataset-anger-5.truncation.signal_end_over_limit; dataset-fear-0.truncation.signal_end_over_limit; dataset-fear-1.truncation.signal_end_over_limit; dataset-fear-2.truncation.signal_end_over_limit; dataset-fear-3.truncation.signal_end_over_limit; dataset-fear-4.truncation.signal_end_over_limit; dataset-fear-5.truncation.signal_end_over_limit; dataset-joy-0.truncation.signal_end_over_limit; dataset-joy-1.truncation.signal_end_over_limit; dataset-joy-2.truncation.signal_end_over_limit; dataset-joy-3.truncation.signal_end_over_limit; dataset-joy-4.truncation.signal_end_over_limit; dataset-joy-5.truncation.signal_end_over_limit; dataset-sadness-0.truncation.signal_end_over_limit; dataset-sadness-1.truncation.signal_end_over_limit; dataset-sadness-2.truncation.signal_end_over_limit; dataset-sadness-3.truncation.signal_end_over_limit; dataset-sadness-4.truncation.signal_end_over_limit; dataset-sadness-5.truncation.signal_end_over_limit; dataset-surprise-0.truncation.signal_end_over_limit; dataset-surprise-1.truncation.signal_end_over_limit; dataset-surprise-2.truncation.signal_end_over_limit; dataset-surprise-3.truncation.signal_end_over_limit; dataset-surprise-4.truncation.signal_end_over_limit; dataset-surprise-5.truncation.signal_end_over_limit; handwritten-disgust-0.truncation.signal_end_over_limit; handwritten-disgust-1.truncation.signal_end_over_limit; handwritten-disgust-2.truncation.signal_end_over_limit; handwritten-disgust-3.truncation.signal_end_over_limit; handwritten-disgust-4.truncation.signal_end_over_limit; handwritten-disgust-5.truncation.signal_end_over_limit; handwritten-neutral-0.truncation.signal_end_over_limit; handwritten-neutral-1.truncation.signal_end_over_limit; handwritten-neutral-2.truncation.signal_end_over_limit; handwritten-neutral-3.truncation.signal_end_over_limit; handwritten-neutral-4.truncation.signal_end_over_limit; handwritten-neutral-5.truncation.signal_end_over_limit;

V1 → V2 evidence · 0 remaining / 42 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 42 cases

dataset-anger-0.truncation.signal_end_over_limit; dataset-anger-1.truncation.signal_end_over_limit; dataset-anger-2.truncation.signal_end_over_limit; dataset-anger-3.truncation.signal_end_over_limit; dataset-anger-4.truncation.signal_end_over_limit; dataset-anger-5.truncation.signal_end_over_limit; dataset-fear-0.truncation.signal_end_over_limit; dataset-fear-1.truncation.signal_end_over_limit; dataset-fear-2.truncation.signal_end_over_limit; dataset-fear-3.truncation.signal_end_over_limit; dataset-fear-4.truncation.signal_end_over_limit; dataset-fear-5.truncation.signal_end_over_limit; dataset-joy-0.truncation.signal_end_over_limit; dataset-joy-1.truncation.signal_end_over_limit; dataset-joy-2.truncation.signal_end_over_limit; dataset-joy-3.truncation.signal_end_over_limit; dataset-joy-4.truncation.signal_end_over_limit; dataset-joy-5.truncation.signal_end_over_limit; dataset-sadness-0.truncation.signal_end_over_limit; dataset-sadness-1.truncation.signal_end_over_limit; dataset-sadness-2.truncation.signal_end_over_limit; dataset-sadness-3.truncation.signal_end_over_limit; dataset-sadness-4.truncation.signal_end_over_limit; dataset-sadness-5.truncation.signal_end_over_limit; dataset-surprise-0.truncation.signal_end_over_limit; dataset-surprise-1.truncation.signal_end_over_limit; dataset-surprise-2.truncation.signal_end_over_limit; dataset-surprise-3.truncation.signal_end_over_limit; dataset-surprise-4.truncation.signal_end_over_limit; dataset-surprise-5.truncation.signal_end_over_limit; handwritten-disgust-0.truncation.signal_end_over_limit; handwritten-disgust-1.truncation.signal_end_over_limit; handwritten-disgust-2.truncation.signal_end_over_limit; handwritten-disgust-3.truncation.signal_end_over_limit; handwritten-disgust-4.truncation.signal_end_over_limit; handwritten-disgust-5.truncation.signal_end_over_limit; handwritten-neutral-0.truncation.signal_end_over_limit; handwritten-neutral-1.truncation.signal_end_over_limit; handwritten-neutral-2.truncation.signal_end_over_limit; handwritten-neutral-3.truncation.signal_end_over_limit; handwritten-neutral-4.truncation.signal_end_over_limit; handwritten-neutral-5.truncation.signal_end_over_limit;

truncation/signal_start_over_limit: unhandled server error

Truncation / signal_start_over_limit · unhandled 5xx · 42 affected cases · GOLD

42 case(s) in truncation/signal_start_over_limit produced unhandled server error -- the endpoint returned an unhandled 5xx; where health still answered afterwards the request failed but the service survived, which is not unavailability. Severity is scored from the single highest-scoring instance; the 42 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Add a scoped exception handler for this failure mode that returns a typed error response instead of an unhandled traceback.

+ Supporting attack IDs · V1 · 42 cases

dataset-anger-0.truncation.signal_start_over_limit; dataset-anger-1.truncation.signal_start_over_limit; dataset-anger-2.truncation.signal_start_over_limit; dataset-anger-3.truncation.signal_start_over_limit; dataset-anger-4.truncation.signal_start_over_limit; dataset-anger-5.truncation.signal_start_over_limit; dataset-fear-0.truncation.signal_start_over_limit; dataset-fear-1.truncation.signal_start_over_limit; dataset-fear-2.truncation.signal_start_over_limit; dataset-fear-3.truncation.signal_start_over_limit; dataset-fear-4.truncation.signal_start_over_limit; dataset-fear-5.truncation.signal_start_over_limit; dataset-joy-0.truncation.signal_start_over_limit; dataset-joy-1.truncation.signal_start_over_limit; dataset-joy-2.truncation.signal_start_over_limit; dataset-joy-3.truncation.signal_start_over_limit; dataset-joy-4.truncation.signal_start_over_limit; dataset-joy-5.truncation.signal_start_over_limit; dataset-sadness-0.truncation.signal_start_over_limit; dataset-sadness-1.truncation.signal_start_over_limit; dataset-sadness-2.truncation.signal_start_over_limit; dataset-sadness-3.truncation.signal_start_over_limit; dataset-sadness-4.truncation.signal_start_over_limit; dataset-sadness-5.truncation.signal_start_over_limit; dataset-surprise-0.truncation.signal_start_over_limit; dataset-surprise-1.truncation.signal_start_over_limit; dataset-surprise-2.truncation.signal_start_over_limit; dataset-surprise-3.truncation.signal_start_over_limit; dataset-surprise-4.truncation.signal_start_over_limit; dataset-surprise-5.truncation.signal_start_over_limit; handwritten-disgust-0.truncation.signal_start_over_limit; handwritten-disgust-1.truncation.signal_start_over_limit; handwritten-disgust-2.truncation.signal_start_over_limit; handwritten-disgust-3.truncation.signal_start_over_limit; handwritten-disgust-4.truncation.signal_start_over_limit; handwritten-disgust-5.truncation.signal_start_over_limit; handwritten-neutral-0.truncation.signal_start_over_limit; handwritten-neutral-1.truncation.signal_start_over_limit; handwritten-neutral-2.truncation.signal_start_over_limit; handwritten-neutral-3.truncation.signal_start_over_limit; handwritten-neutral-4.truncation.signal_start_over_limit; handwritten-neutral-5.truncation.signal_start_over_limit;

V1 → V2 evidence · 0 remaining / 42 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 42 cases

dataset-anger-0.truncation.signal_start_over_limit; dataset-anger-1.truncation.signal_start_over_limit; dataset-anger-2.truncation.signal_start_over_limit; dataset-anger-3.truncation.signal_start_over_limit; dataset-anger-4.truncation.signal_start_over_limit; dataset-anger-5.truncation.signal_start_over_limit; dataset-fear-0.truncation.signal_start_over_limit; dataset-fear-1.truncation.signal_start_over_limit; dataset-fear-2.truncation.signal_start_over_limit; dataset-fear-3.truncation.signal_start_over_limit; dataset-fear-4.truncation.signal_start_over_limit; dataset-fear-5.truncation.signal_start_over_limit; dataset-joy-0.truncation.signal_start_over_limit; dataset-joy-1.truncation.signal_start_over_limit; dataset-joy-2.truncation.signal_start_over_limit; dataset-joy-3.truncation.signal_start_over_limit; dataset-joy-4.truncation.signal_start_over_limit; dataset-joy-5.truncation.signal_start_over_limit; dataset-sadness-0.truncation.signal_start_over_limit; dataset-sadness-1.truncation.signal_start_over_limit; dataset-sadness-2.truncation.signal_start_over_limit; dataset-sadness-3.truncation.signal_start_over_limit; dataset-sadness-4.truncation.signal_start_over_limit; dataset-sadness-5.truncation.signal_start_over_limit; dataset-surprise-0.truncation.signal_start_over_limit; dataset-surprise-1.truncation.signal_start_over_limit; dataset-surprise-2.truncation.signal_start_over_limit; dataset-surprise-3.truncation.signal_start_over_limit; dataset-surprise-4.truncation.signal_start_over_limit; dataset-surprise-5.truncation.signal_start_over_limit; handwritten-disgust-0.truncation.signal_start_over_limit; handwritten-disgust-1.truncation.signal_start_over_limit; handwritten-disgust-2.truncation.signal_start_over_limit; handwritten-disgust-3.truncation.signal_start_over_limit; handwritten-disgust-4.truncation.signal_start_over_limit; handwritten-disgust-5.truncation.signal_start_over_limit; handwritten-neutral-0.truncation.signal_start_over_limit; handwritten-neutral-1.truncation.signal_start_over_limit; handwritten-neutral-2.truncation.signal_start_over_limit; handwritten-neutral-3.truncation.signal_start_over_limit; handwritten-neutral-4.truncation.signal_start_over_limit; handwritten-neutral-5.truncation.signal_start_over_limit;

encoding/compatibility_fullwidth: prediction flip

Encoding / compatibility_fullwidth · prediction flip · 29 affected cases · SILVER

29 case(s) in encoding/compatibility_fullwidth produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 29 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 29 cases

dataset-anger-0.encoding.compatibility_fullwidth.whole; dataset-anger-1.encoding.compatibility_fullwidth.whole; dataset-anger-2.encoding.compatibility_fullwidth.whole; dataset-anger-3.encoding.compatibility_fullwidth.whole; dataset-anger-4.encoding.compatibility_fullwidth.whole; dataset-fear-0.encoding.compatibility_fullwidth.whole; dataset-fear-1.encoding.compatibility_fullwidth.whole; dataset-fear-2.encoding.compatibility_fullwidth.whole; dataset-fear-3.encoding.compatibility_fullwidth.whole; dataset-fear-4.encoding.compatibility_fullwidth.whole; dataset-fear-5.encoding.compatibility_fullwidth.whole; dataset-joy-0.encoding.compatibility_fullwidth.whole; dataset-joy-2.encoding.compatibility_fullwidth.whole; dataset-joy-3.encoding.compatibility_fullwidth.whole; dataset-joy-4.encoding.compatibility_fullwidth.whole; dataset-joy-5.encoding.compatibility_fullwidth.whole; dataset-sadness-0.encoding.compatibility_fullwidth.whole; dataset-sadness-1.encoding.compatibility_fullwidth.whole; dataset-sadness-2.encoding.compatibility_fullwidth.whole; dataset-sadness-3.encoding.compatibility_fullwidth.whole; dataset-sadness-4.encoding.compatibility_fullwidth.whole; dataset-sadness-5.encoding.compatibility_fullwidth.whole; dataset-surprise-1.encoding.compatibility_fullwidth.whole; dataset-surprise-2.encoding.compatibility_fullwidth.whole; dataset-surprise-4.encoding.compatibility_fullwidth.whole; handwritten-disgust-1.encoding.compatibility_fullwidth.whole; handwritten-disgust-3.encoding.compatibility_fullwidth.whole; handwritten-disgust-5.encoding.compatibility_fullwidth.whole;

V1 → V2 evidence · 0 remaining / 29 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 29 cases

dataset-anger-0.encoding.compatibility_fullwidth.whole; dataset-anger-1.encoding.compatibility_fullwidth.whole; dataset-anger-2.encoding.compatibility_fullwidth.whole; dataset-anger-3.encoding.compatibility_fullwidth.whole; dataset-anger-4.encoding.compatibility_fullwidth.whole; dataset-fear-0.encoding.compatibility_fullwidth.whole; dataset-fear-1.encoding.compatibility_fullwidth.whole; dataset-fear-2.encoding.compatibility_fullwidth.whole; dataset-fear-3.encoding.compatibility_fullwidth.whole; dataset-fear-4.encoding.compatibility_fullwidth.whole; dataset-fear-5.encoding.compatibility_fullwidth.whole; dataset-joy-0.encoding.compatibility_fullwidth.whole; dataset-joy-2.encoding.compatibility_fullwidth.whole; dataset-joy-3.encoding.compatibility_fullwidth.whole; dataset-joy-4.encoding.compatibility_fullwidth.whole; dataset-joy-5.encoding.compatibility_fullwidth.whole; dataset-sadness-0.encoding.compatibility_fullwidth.whole; dataset-sadness-1.encoding.compatibility_fullwidth.whole; dataset-sadness-2.encoding.compatibility_fullwidth.whole; dataset-sadness-3.encoding.compatibility_fullwidth.whole; dataset-sadness-4.encoding.compatibility_fullwidth.whole; dataset-sadness-5.encoding.compatibility_fullwidth.whole; dataset-surprise-1.encoding.compatibility_fullwidth.whole; dataset-surprise-2.encoding.compatibility_fullwidth.whole; dataset-surprise-4.encoding.compatibility_fullwidth.whole; handwritten-disgust-0.encoding.compatibility_fullwidth.whole; handwritten-disgust-1.encoding.compatibility_fullwidth.whole; handwritten-disgust-3.encoding.compatibility_fullwidth.whole; handwritten-disgust-5.encoding.compatibility_fullwidth.whole;

encoding/homoglyph: prediction flip

Encoding / homoglyph · prediction flip · 37 affected cases · SILVER

37 case(s) in encoding/homoglyph produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 37 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 37 cases

dataset-anger-1.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.trailing.d3; dataset-anger-2.encoding.homoglyph.trailing.d5; dataset-anger-3.encoding.homoglyph.trailing.d1; dataset-anger-3.encoding.homoglyph.trailing.d3; dataset-anger-3.encoding.homoglyph.trailing.d5; dataset-fear-0.encoding.homoglyph.leading.d5; dataset-fear-0.encoding.homoglyph.middle.d1; dataset-fear-0.encoding.homoglyph.middle.d3; dataset-fear-0.encoding.homoglyph.middle.d5; dataset-fear-0.encoding.homoglyph.trailing.d5; dataset-fear-4.encoding.homoglyph.leading.d5; dataset-joy-0.encoding.homoglyph.leading.d1; dataset-joy-0.encoding.homoglyph.leading.d3; dataset-joy-0.encoding.homoglyph.leading.d5; dataset-joy-3.encoding.homoglyph.trailing.d3; dataset-joy-3.encoding.homoglyph.trailing.d5; dataset-joy-5.encoding.homoglyph.trailing.d5; dataset-sadness-0.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.leading.d5; dataset-sadness-3.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.trailing.d5; dataset-sadness-4.encoding.homoglyph.leading.d5; dataset-sadness-5.encoding.homoglyph.leading.d3; dataset-sadness-5.encoding.homoglyph.leading.d5; dataset-surprise-2.encoding.homoglyph.trailing.d1; dataset-surprise-2.encoding.homoglyph.trailing.d3; dataset-surprise-2.encoding.homoglyph.trailing.d5; handwritten-neutral-1.encoding.homoglyph.middle.d3; handwritten-neutral-1.encoding.homoglyph.middle.d5; handwritten-neutral-2.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.middle.d1; handwritten-neutral-5.encoding.homoglyph.middle.d3; handwritten-neutral-5.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.trailing.d1; handwritten-neutral-5.encoding.homoglyph.trailing.d5;

V1 → V2 evidence · 37 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 37 cases

dataset-anger-1.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.trailing.d3; dataset-anger-2.encoding.homoglyph.trailing.d5; dataset-anger-3.encoding.homoglyph.trailing.d1; dataset-anger-3.encoding.homoglyph.trailing.d3; dataset-anger-3.encoding.homoglyph.trailing.d5; dataset-fear-0.encoding.homoglyph.leading.d5; dataset-fear-0.encoding.homoglyph.middle.d1; dataset-fear-0.encoding.homoglyph.middle.d3; dataset-fear-0.encoding.homoglyph.middle.d5; dataset-fear-0.encoding.homoglyph.trailing.d5; dataset-fear-4.encoding.homoglyph.leading.d5; dataset-joy-0.encoding.homoglyph.leading.d1; dataset-joy-0.encoding.homoglyph.leading.d3; dataset-joy-0.encoding.homoglyph.leading.d5; dataset-joy-3.encoding.homoglyph.trailing.d3; dataset-joy-3.encoding.homoglyph.trailing.d5; dataset-joy-5.encoding.homoglyph.trailing.d5; dataset-sadness-0.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.leading.d5; dataset-sadness-3.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.trailing.d5; dataset-sadness-4.encoding.homoglyph.leading.d5; dataset-sadness-5.encoding.homoglyph.leading.d3; dataset-sadness-5.encoding.homoglyph.leading.d5; dataset-surprise-2.encoding.homoglyph.trailing.d1; dataset-surprise-2.encoding.homoglyph.trailing.d3; dataset-surprise-2.encoding.homoglyph.trailing.d5; handwritten-neutral-1.encoding.homoglyph.middle.d3; handwritten-neutral-1.encoding.homoglyph.middle.d5; handwritten-neutral-2.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.middle.d1; handwritten-neutral-5.encoding.homoglyph.middle.d3; handwritten-neutral-5.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.trailing.d1; handwritten-neutral-5.encoding.homoglyph.trailing.d5;

encoding/homoglyph_greek: prediction flip

Encoding / homoglyph_greek · prediction flip · 21 affected cases · SILVER

21 case(s) in encoding/homoglyph_greek produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 21 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 21 cases

dataset-anger-1.encoding.homoglyph_greek.leading.d1; dataset-anger-1.encoding.homoglyph_greek.leading.d3; dataset-anger-3.encoding.homoglyph_greek.middle.d3; dataset-anger-3.encoding.homoglyph_greek.trailing.d1; dataset-anger-3.encoding.homoglyph_greek.trailing.d3; dataset-fear-0.encoding.homoglyph_greek.leading.d1; dataset-fear-0.encoding.homoglyph_greek.leading.d3; dataset-fear-0.encoding.homoglyph_greek.middle.d3; dataset-fear-0.encoding.homoglyph_greek.trailing.d3; dataset-fear-3.encoding.homoglyph_greek.leading.d3; dataset-fear-4.encoding.homoglyph_greek.leading.d1; dataset-fear-4.encoding.homoglyph_greek.leading.d3; dataset-joy-0.encoding.homoglyph_greek.leading.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d3; dataset-sadness-5.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.leading.d1; handwritten-neutral-2.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3; handwritten-neutral-3.encoding.homoglyph_greek.leading.d3; handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3;

V1 → V2 evidence · 21 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 21 cases

dataset-anger-1.encoding.homoglyph_greek.leading.d1; dataset-anger-1.encoding.homoglyph_greek.leading.d3; dataset-anger-3.encoding.homoglyph_greek.middle.d3; dataset-anger-3.encoding.homoglyph_greek.trailing.d1; dataset-anger-3.encoding.homoglyph_greek.trailing.d3; dataset-fear-0.encoding.homoglyph_greek.leading.d1; dataset-fear-0.encoding.homoglyph_greek.leading.d3; dataset-fear-0.encoding.homoglyph_greek.middle.d3; dataset-fear-0.encoding.homoglyph_greek.trailing.d3; dataset-fear-3.encoding.homoglyph_greek.leading.d3; dataset-fear-4.encoding.homoglyph_greek.leading.d1; dataset-fear-4.encoding.homoglyph_greek.leading.d3; dataset-joy-0.encoding.homoglyph_greek.leading.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d3; dataset-sadness-5.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.leading.d1; handwritten-neutral-2.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3; handwritten-neutral-3.encoding.homoglyph_greek.leading.d3; handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3;

F-v1-013 **Medium · 59.8/100** Remaining

perturbation/capitalization: prediction flip

Perturbation / capitalization · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/capitalization produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 7 cases

dataset-anger-1.perturbation.capitalization; dataset-anger-3.perturbation.capitalization; dataset-fear-0.perturbation.capitalization; dataset-fear-3.perturbation.capitalization; dataset-joy-0.perturbation.capitalization; dataset-sadness-3.perturbation.capitalization; dataset-sadness-4.perturbation.capitalization;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.capitalization; dataset-anger-3.perturbation.capitalization; dataset-fear-0.perturbation.capitalization; dataset-fear-3.perturbation.capitalization; dataset-joy-0.perturbation.capitalization; dataset-sadness-3.perturbation.capitalization; dataset-sadness-4.perturbation.capitalization;

F-v1-014

Medium · 59.8/100

Remaining

perturbation/delete_char: prediction flip

Perturbation / delete_char · prediction flip · 5 affected cases · SILVER

5 case(s) in perturbation/delete_char produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 5 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 5 cases

dataset-anger-1.perturbation.delete_char; dataset-anger-2.perturbation.delete_char; dataset-anger-3.perturbation.delete_char; dataset-fear-0.perturbation.delete_char; dataset-fear-3.perturbation.delete_char;

V1 → V2 evidence · 5 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.perturbation.delete_char; dataset-anger-2.perturbation.delete_char; dataset-anger-3.perturbation.delete_char; dataset-fear-0.perturbation.delete_char; dataset-fear-3.perturbation.delete_char;

F-v1-015

Medium · 59.8/100

Remaining

perturbation/insert_char: prediction flip

Perturbation / insert_char · prediction flip · 6 affected cases · SILVER

6 case(s) in perturbation/insert_char produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 6 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 6 cases

dataset-anger-1.perturbation.insert_char; dataset-anger-2.perturbation.insert_char; dataset-anger-3.perturbation.insert_char; dataset-fear-0.perturbation.insert_char; dataset-fear-3.perturbation.insert_char; dataset-joy-0.perturbation.insert_char;

V1 → V2 evidence · 6 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 6 cases

dataset-anger-1.perturbation.insert_char; dataset-anger-2.perturbation.insert_char; dataset-anger-3.perturbation.insert_char; dataset-fear-0.perturbation.insert_char; dataset-fear-3.perturbation.insert_char; dataset-joy-0.perturbation.insert_char;

F-v1-016 **Medium · 59.8/100** Remaining

perturbation/keyboard_typo: prediction flip

Perturbation / keyboard_typo · prediction flip · 5 affected cases · SILVER

5 case(s) in perturbation/keyboard_typo produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 5 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 5 cases

dataset-anger-1.perturbation.keyboard_typo; dataset-anger-3.perturbation.keyboard_typo; dataset-fear-0.perturbation.keyboard_typo; dataset-fear-3.perturbation.keyboard_typo; dataset-joy-0.perturbation.keyboard_typo;

V1 → V2 evidence · 5 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.perturbation.keyboard_typo; dataset-anger-3.perturbation.keyboard_typo; dataset-fear-0.perturbation.keyboard_typo; dataset-fear-3.perturbation.keyboard_typo; dataset-joy-0.perturbation.keyboard_typo;

F-v1-019 **Medium · 59.8/100** Remaining

perturbation/swap_adjacent: prediction flip

Perturbation / swap_adjacent · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/swap_adjacent produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 7 cases

dataset-anger-1.perturbation.swap_adjacent; dataset-anger-2.perturbation.swap_adjacent; dataset-anger-3.perturbation.swap_adjacent; dataset-fear-0.perturbation.swap_adjacent; dataset-fear-3.perturbation.swap_adjacent; dataset-joy-0.perturbation.swap_adjacent; dataset-joy-4.perturbation.swap_adjacent;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.swap_adjacent; dataset-anger-2.perturbation.swap_adjacent; dataset-anger-3.perturbation.swap_adjacent; dataset-fear-0.perturbation.swap_adjacent; dataset-fear-3.perturbation.swap_adjacent; dataset-joy-0.perturbation.swap_adjacent; dataset-joy-4.perturbation.swap_adjacent;

F-v1-020 **Medium · 59.8/100** **Remaining**

perturbation/word_split: prediction flip

Perturbation / word_split · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/word_split produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 7 cases

dataset-anger-1.perturbation.word_split; dataset-anger-2.perturbation.word_split; dataset-anger-3.perturbation.word_split; dataset-fear-0.perturbation.word_split; dataset-fear-3.perturbation.word_split; dataset-joy-0.perturbation.word_split; dataset-joy-4.perturbation.word_split;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.word_split; dataset-anger-2.perturbation.word_split; dataset-anger-3.perturbation.word_split; dataset-fear-0.perturbation.word_split; dataset-fear-3.perturbation.word_split; dataset-joy-0.perturbation.word_split; dataset-joy-4.perturbation.word_split;

F-v1-028 **Medium · 59.8/100** **Resolved**

whitespace/tabs: prediction flip

Whitespace / tabs · prediction flip · 18 affected cases · SILVER

18 case(s) in whitespace/tabs produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 18 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 18 cases

dataset-anger-1.whitespace.tabs; dataset-anger-2.whitespace.tabs; dataset-anger-5.whitespace.tabs; dataset-fear-0.whitespace.tabs; dataset-fear-1.whitespace.tabs; dataset-fear-3.whitespace.tabs; dataset-joy-1.whitespace.tabs; dataset-joy-3.whitespace.tabs; dataset-joy-5.whitespace.tabs; dataset-sadness-5.whitespace.tabs; dataset-surprise-1.whitespace.tabs; dataset-surprise-2.whitespace.tabs; handwritten-disgust-0.whitespace.tabs; handwritten-disgust-1.whitespace.tabs; handwritten-disgust-3.whitespace.tabs; handwritten-disgust-5.whitespace.tabs; handwritten-neutral-3.whitespace.tabs; handwritten-neutral-5.whitespace.tabs;

V1 → V2 evidence · 0 remaining / 18 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 18 cases

dataset-anger-1.whitespace.tabs; dataset-anger-2.whitespace.tabs; dataset-anger-5.whitespace.tabs; dataset-fear-0.whitespace.tabs; dataset-fear-1.whitespace.tabs; dataset-fear-3.whitespace.tabs; dataset-joy-1.whitespace.tabs; dataset-joy-3.whitespace.tabs; dataset-joy-5.whitespace.tabs; dataset-sadness-5.whitespace.tabs; dataset-surprise-1.whitespace.tabs; dataset-surprise-2.whitespace.tabs; handwritten-disgust-0.whitespace.tabs; handwritten-disgust-1.whitespace.tabs; handwritten-disgust-3.whitespace.tabs; handwritten-disgust-5.whitespace.tabs; handwritten-neutral-3.whitespace.tabs; handwritten-neutral-5.whitespace.tabs;

F-v1-025 **Medium · 59.8/100** **Resolved**

whitespace/interior_runs: prediction flip

Whitespace / interior_runs · prediction flip · 19 affected cases · SILVER

19 case(s) in whitespace/interior_runs produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 19 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 19 cases

dataset-anger-0.whitespace.interior_runs; dataset-anger-1.whitespace.interior_runs; dataset-anger-2.whitespace.interior_runs; dataset-anger-4.whitespace.interior_runs; dataset-anger-5.whitespace.interior_runs; dataset-fear-3.whitespace.interior_runs; dataset-fear-4.whitespace.interior_runs; dataset-sadness-0.whitespace.interior_runs; dataset-surprise-2.whitespace.interior_runs; dataset-surprise-3.whitespace.interior_runs; dataset-surprise-4.whitespace.interior_runs; handwritten-disgust-0.whitespace.interior_runs; handwritten-disgust-1.whitespace.interior_runs; handwritten-disgust-3.whitespace.interior_runs; handwritten-disgust-5.whitespace.interior_runs; handwritten-neutral-1.whitespace.interior_runs; handwritten-neutral-2.whitespace.interior_runs; handwritten-neutral-3.whitespace.interior_runs; handwritten-neutral-5.whitespace.interior_runs;

V1 → V2 evidence · 0 remaining / 19 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 19 cases

dataset-anger-0.whitespace.interior_runs; dataset-anger-1.whitespace.interior_runs; dataset-anger-2.whitespace.interior_runs; dataset-anger-4.whitespace.interior_runs; dataset-anger-5.whitespace.interior_runs; dataset-fear-3.whitespace.interior_runs; dataset-fear-4.whitespace.interior_runs; dataset-sadness-0.whitespace.interior_runs; dataset-surprise-2.whitespace.interior_runs; dataset-surprise-3.whitespace.interior_runs; dataset-surprise-4.whitespace.interior_runs; handwritten-disgust-0.whitespace.interior_runs; handwritten-disgust-1.whitespace.interior_runs; handwritten-disgust-3.whitespace.interior_runs; handwritten-disgust-5.whitespace.interior_runs; handwritten-neutral-1.whitespace.interior_runs; handwritten-neutral-2.whitespace.interior_runs; handwritten-neutral-3.whitespace.interior_runs; handwritten-neutral-5.whitespace.interior_runs;

perturbation/keyboard_typo_second_word: prediction flip

Perturbation / keyboard_typo_second_word · prediction flip · 2 affected cases · SILVER

2 case(s) in perturbation/keyboard_typo_second_word produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 2 cases

dataset-sadness-0.perturbation.keyboard_typo_second_word; dataset-sadness-5.perturbation.keyboard_typo_second_word;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-sadness-0.perturbation.keyboard_typo_second_word; dataset-sadness-5.perturbation.keyboard_typo_second_word;

encoding/deletion_backspace: prediction flip

Encoding / deletion_backspace · prediction flip · 7 affected cases · SILVER

7 case(s) in encoding/deletion_backspace produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 7 cases

dataset-anger-1.encoding.deletion_backspace.leading.d1; dataset-anger-3.encoding.deletion_backspace.trailing.d1; dataset-joy-0.encoding.deletion_backspace.leading.d1; dataset-sadness-5.encoding.deletion_backspace.leading.d1; dataset-sadness-5.encoding.deletion_backspace.trailing.d1; dataset-surprise-2.encoding.deletion_backspace.trailing.d1; handwritten-neutral-5.encoding.deletion_backspace.middle.d1;

V1 → V2 evidence · 5 remaining / 2 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.encoding.deletion_backspace.leading.d1; dataset-anger-3.encoding.deletion_backspace.trailing.d1; dataset-joy-0.encoding.deletion_backspace.leading.d1; dataset-sadness-5.encoding.deletion_backspace.leading.d1; dataset-surprise-2.encoding.deletion_backspace.trailing.d1;

+ Improved original cases · 2 cases

dataset-sadness-5.encoding.deletion_backspace.trailing.d1; handwritten-neutral-5.encoding.deletion_backspace.middle.d1;

F-v1-007

Medium · 59.7/100

Resolved

encoding/invisible_zero_width: prediction flip

Encoding / invisible_zero_width · prediction flip · 7 affected cases · SILVER

7 case(s) in encoding/invisible_zero_width produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 7 cases

dataset-anger-1.encoding.invisible_zero_width.leading.d3; dataset-anger-3.encoding.invisible_zero_width.trailing.d3; dataset-sadness-5.encoding.invisible_zero_width.leading.d3; dataset-surprise-2.encoding.invisible_zero_width.trailing.d3; handwritten-neutral-2.encoding.invisible_zero_width.leading.d3; handwritten-neutral-5.encoding.invisible_zero_width.middle.d3; handwritten-neutral-5.encoding.invisible_zero_width.trailing.d3;

V1 → V2 evidence · 0 remaining / 7 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 7 cases

dataset-anger-1.encoding.invisible_zero_width.leading.d3; dataset-anger-3.encoding.invisible_zero_width.trailing.d3; dataset-sadness-5.encoding.invisible_zero_width.leading.d3; dataset-surprise-2.encoding.invisible_zero_width.trailing.d3; handwritten-neutral-2.encoding.invisible_zero_width.leading.d3; handwritten-neutral-5.encoding.invisible_zero_width.middle.d3; handwritten-neutral-5.encoding.invisible_zero_width.trailing.d3;

F-v1-018

Medium · 59.7/100

Remaining

perturbation/misspell: prediction flip

Perturbation / misspell · prediction flip · 6 affected cases · SILVER

6 case(s) in perturbation/misspell produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 6 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 6 cases

dataset-anger-1.perturbation.misspell; dataset-anger-2.perturbation.misspell; dataset-anger-3.perturbation.misspell; dataset-fear-3.perturbation.misspell; dataset-joy-0.perturbation.misspell; dataset-joy-4.perturbation.misspell;

V1 → V2 evidence · 6 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 6 cases

dataset-anger-1.perturbation.misspell; dataset-anger-2.perturbation.misspell; dataset-anger-3.perturbation.misspell; dataset-fear-3.perturbation.misspell; dataset-joy-0.perturbation.misspell; dataset-joy-4.perturbation.misspell;

F-v1-027

Medium · 59.7/100

Resolved

whitespace/newlines: prediction flip

Whitespace / newlines · prediction flip · 13 affected cases · SILVER

13 case(s) in whitespace/newlines produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 13 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 13 cases

dataset-anger-1.whitespace.newlines; dataset-anger-2.whitespace.newlines; dataset-anger-5.whitespace.newlines; dataset-fear-1.whitespace.newlines; dataset-fear-3.whitespace.newlines; dataset-joy-1.whitespace.newlines; dataset-joy-3.whitespace.newlines; dataset-sadness-5.whitespace.newlines; dataset-surprise-1.whitespace.newlines; dataset-surprise-2.whitespace.newlines; dataset-surprise-3.whitespace.newlines; handwritten-neutral-3.whitespace.newlines; handwritten-neutral-5.whitespace.newlines;

V1 → V2 evidence · 0 remaining / 13 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 13 cases

dataset-anger-1.whitespace.newlines; dataset-anger-2.whitespace.newlines; dataset-anger-5.whitespace.newlines; dataset-fear-1.whitespace.newlines; dataset-fear-3.whitespace.newlines; dataset-joy-1.whitespace.newlines; dataset-joy-3.whitespace.newlines; dataset-sadness-5.whitespace.newlines; dataset-surprise-1.whitespace.newlines; dataset-surprise-2.whitespace.newlines; dataset-surprise-3.whitespace.newlines; handwritten-neutral-3.whitespace.newlines; handwritten-neutral-5.whitespace.newlines;

F-v1-029

Medium · 59.7/100

Resolved

whitespace/unicode_spaces: prediction flip

Whitespace / unicode_spaces · prediction flip · 10 affected cases · SILVER

10 case(s) in whitespace/unicode_spaces produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 10 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 10 cases

dataset-anger-1.whitespace.unicode_spaces; dataset-anger-2.whitespace.unicode_spaces; dataset-fear-1.whitespace.unicode_spaces; dataset-fear-3.whitespace.unicode_spaces; dataset-joy-1.whitespace.unicode_spaces; dataset-surprise-1.whitespace.unicode_spaces; dataset-surprise-2.whitespace.unicode_spaces; handwritten-disgust-0.whitespace.unicode_spaces; handwritten-disgust-1.whitespace.unicode_spaces; handwritten-disgust-5.whitespace.unicode_spaces;

V1 → V2 evidence · 0 remaining / 10 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 10 cases

dataset-anger-1.whitespace.unicode_spaces; dataset-anger-2.whitespace.unicode_spaces; dataset-fear-1.whitespace.unicode_spaces; dataset-fear-3.whitespace.unicode_spaces; dataset-joy-1.whitespace.unicode_spaces; dataset-surprise-1.whitespace.unicode_spaces; dataset-surprise-2.whitespace.unicode_spaces; handwritten-disgust-0.whitespace.unicode_spaces; handwritten-disgust-1.whitespace.unicode_spaces; handwritten-disgust-5.whitespace.unicode_spaces;

F-v1-026

Medium · 56.5/100

Resolved

whitespace/leading_trailing: prediction flip

Whitespace / leading_trailing · prediction flip · 8 affected cases · SILVER

8 case(s) in whitespace/leading_trailing produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 8 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 8 cases

dataset-fear-1.whitespace.leading_trailing; dataset-joy-0.whitespace.leading_trailing; dataset-sadness-5.whitespace.leading_trailing; dataset-surprise-3.whitespace.leading_trailing; handwritten-neutral-1.whitespace.leading_trailing; handwritten-neutral-2.whitespace.leading_trailing; handwritten-neutral-3.whitespace.leading_trailing; handwritten-neutral-5.whitespace.leading_trailing;

V1 → V2 evidence · 0 remaining / 8 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 8 cases

dataset-fear-1.whitespace.leading_trailing; dataset-joy-0.whitespace.leading_trailing; dataset-sadness-5.whitespace.leading_trailing; dataset-surprise-3.whitespace.leading_trailing; handwritten-neutral-1.whitespace.leading_trailing; handwritten-neutral-2.whitespace.leading_trailing; handwritten-neutral-3.whitespace.leading_trailing; handwritten-neutral-5.whitespace.leading_trailing;

F-v1-022 **Medium · 50.5/100** Remaining

truncation/signal_end_short: prediction flip

Truncation / signal_end_short · prediction flip · 2 affected cases · SILVER

2 case(s) in truncation/signal_end_short produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 2 cases

dataset-anger-1.truncation.signal_end_short; dataset-sadness-5.truncation.signal_end_short;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-anger-1.truncation.signal_end_short; dataset-sadness-5.truncation.signal_end_short;

F-v1-024 **Medium · 50.5/100** Remaining

truncation/signal_start_short: prediction flip

Truncation / signal_start_short · prediction flip · 2 affected cases · SILVER

2 case(s) in truncation/signal_start_short produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V1 · 2 cases

dataset-anger-1.truncation.signal_start_short; dataset-sadness-5.truncation.signal_start_short;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-anger-1.truncation.signal_start_short; dataset-sadness-5.truncation.signal_start_short;

F-v1-008 Medium · 40/100 Resolved

malformed/extra_field: invalid input accepted

Malformed / extra_field · invalid input accepted · 1 affected case · GOLD

1 case(s) in malformed/extra_field produced invalid input accepted -- the manifest oracle says this request should have been rejected, but the endpoint accepted it and answered. Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Reject the malformed or out-of-schema input with a 4xx response instead of processing it.

+ Supporting attack IDs · V1 · 1 case

malformed.extra_field;

V1 → V2 evidence · 0 remaining / 1 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Improved original cases · 1 case

malformed.extra_field;

F-v1-010 Low · 33.7/100 Remaining

malformed/oversized_10mb: slow response

Malformed / oversized_10mb · slow response · 1 affected case · GOLD

1 case(s) in malformed/oversized_10mb produced slow response -- the response landed in the slow band (2 seconds up to the timeout). Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Profile and bound the slow code path so typical requests stay under the 2-second latency band.

+ Supporting attack IDs · V1 · 1 case

malformed.oversized_10mb;

V1 → V2 evidence · 1 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 1 case

malformed.oversized_10mb;

V2 finding register 14 groups

F-v2-002 · encoding/homoglyph · prediction flip

F-v2-003 · encoding/homoglyph_greek · prediction flip

F-v2-005 · perturbation/capitalization · prediction flip

F-v2-006 · perturbation/delete_char · prediction flip

F-v2-007 · perturbation/insert_char · prediction flip

F-v2-008 · perturbation/keyboard_typo · prediction flip

F-v2-011 · perturbation/swap_adjacent · prediction flip

F-v2-012 · perturbation/word_split · prediction flip

F-v2-009 · perturbation/keyboard_typo_second_word · prediction flip

F-v2-001 · encoding/deletion_backspace · prediction flip

F-v2-010 · perturbation/misspell · prediction flip

F-v2-013 · truncation/signal_end_short · prediction flip

F-v2-002 **Medium · 59.8/100** Remaining

encoding/homoglyph: prediction flip

Encoding / homoglyph · prediction flip · 37 affected cases · SILVER

37 case(s) in encoding/homoglyph produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 37 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 37 cases

dataset-anger-1.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.trailing.d3; dataset-anger-2.encoding.homoglyph.trailing.d5; dataset-anger-3.encoding.homoglyph.trailing.d1; dataset-anger-3.encoding.homoglyph.trailing.d3; dataset-anger-3.encoding.homoglyph.trailing.d5; dataset-fear-0.encoding.homoglyph.leading.d5; dataset-fear-0.encoding.homoglyph.middle.d1; dataset-fear-0.encoding.homoglyph.middle.d3; dataset-fear-0.encoding.homoglyph.middle.d5; dataset-fear-0.encoding.homoglyph.trailing.d5; dataset-fear-4.encoding.homoglyph.leading.d5; dataset-joy-0.encoding.homoglyph.leading.d1; dataset-joy-0.encoding.homoglyph.leading.d3; dataset-joy-0.encoding.homoglyph.leading.d5; dataset-joy-3.encoding.homoglyph.trailing.d3; dataset-joy-3.encoding.homoglyph.trailing.d5; dataset-joy-5.encoding.homoglyph.trailing.d5; dataset-sadness-0.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.leading.d5; dataset-sadness-3.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.trailing.d5; dataset-sadness-4.encoding.homoglyph.leading.d5; dataset-sadness-5.encoding.homoglyph.leading.d3; dataset-sadness-5.encoding.homoglyph.leading.d5; dataset-surprise-2.encoding.homoglyph.trailing.d1; dataset-surprise-2.encoding.homoglyph.trailing.d3; dataset-surprise-2.encoding.homoglyph.trailing.d5; handwritten-neutral-1.encoding.homoglyph.middle.d3; handwritten-neutral-1.encoding.homoglyph.middle.d5; handwritten-neutral-2.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.middle.d1; handwritten-neutral-5.encoding.homoglyph.middle.d3; handwritten-neutral-5.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.trailing.d1; handwritten-neutral-5.encoding.homoglyph.trailing.d5;

V1 → V2 evidence · 37 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 37 cases

dataset-anger-1.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.leading.d3; dataset-anger-2.encoding.homoglyph.trailing.d3; dataset-anger-2.encoding.homoglyph.trailing.d5; dataset-anger-3.encoding.homoglyph.trailing.d1; dataset-anger-3.encoding.homoglyph.trailing.d3; dataset-anger-3.encoding.homoglyph.trailing.d5; dataset-fear-0.encoding.homoglyph.leading.d5; dataset-fear-0.encoding.homoglyph.middle.d1; dataset-fear-0.encoding.homoglyph.middle.d3; dataset-fear-0.encoding.homoglyph.middle.d5; dataset-fear-0.encoding.homoglyph.trailing.d5; dataset-fear-4.encoding.homoglyph.leading.d5; dataset-joy-0.encoding.homoglyph.leading.d1; dataset-joy-0.encoding.homoglyph.leading.d3; dataset-joy-0.encoding.homoglyph.leading.d5; dataset-joy-3.encoding.homoglyph.trailing.d3; dataset-joy-3.encoding.homoglyph.trailing.d5; dataset-joy-5.encoding.homoglyph.trailing.d5; dataset-sadness-0.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.leading.d5; dataset-sadness-3.encoding.homoglyph.middle.d5; dataset-sadness-3.encoding.homoglyph.trailing.d5; dataset-sadness-4.encoding.homoglyph.leading.d5; dataset-sadness-5.encoding.homoglyph.leading.d3; dataset-sadness-5.encoding.homoglyph.leading.d5; dataset-surprise-2.encoding.homoglyph.trailing.d1; dataset-surprise-2.encoding.homoglyph.trailing.d3; dataset-surprise-2.encoding.homoglyph.trailing.d5; handwritten-neutral-1.encoding.homoglyph.middle.d3; handwritten-neutral-1.encoding.homoglyph.middle.d5; handwritten-neutral-2.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.middle.d1; handwritten-neutral-5.encoding.homoglyph.middle.d3; handwritten-neutral-5.encoding.homoglyph.middle.d5; handwritten-neutral-5.encoding.homoglyph.trailing.d1; handwritten-neutral-5.encoding.homoglyph.trailing.d5;

encoding/homoglyph_greek: prediction flip

Encoding / homoglyph_greek · prediction flip · 21 affected cases · SILVER

21 case(s) in encoding/homoglyph_greek produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 21 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 21 cases

dataset-anger-1.encoding.homoglyph_greek.leading.d1; dataset-anger-1.encoding.homoglyph_greek.leading.d3; dataset-anger-3.encoding.homoglyph_greek.middle.d3; dataset-anger-3.encoding.homoglyph_greek.trailing.d1; dataset-anger-3.encoding.homoglyph_greek.trailing.d3; dataset-fear-0.encoding.homoglyph_greek.leading.d1; dataset-fear-0.encoding.homoglyph_greek.leading.d3; dataset-fear-0.encoding.homoglyph_greek.middle.d3; dataset-fear-0.encoding.homoglyph_greek.trailing.d3; dataset-fear-3.encoding.homoglyph_greek.leading.d3; dataset-fear-4.encoding.homoglyph_greek.leading.d1; dataset-fear-4.encoding.homoglyph_greek.leading.d3; dataset-joy-0.encoding.homoglyph_greek.leading.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d3; dataset-sadness-5.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.leading.d1; handwritten-neutral-2.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3; handwritten-neutral-3.encoding.homoglyph_greek.leading.d3; handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3;

V1 → V2 evidence · 21 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 21 cases

dataset-anger-1.encoding.homoglyph_greek.leading.d1; dataset-anger-1.encoding.homoglyph_greek.leading.d3; dataset-anger-3.encoding.homoglyph_greek.middle.d3; dataset-anger-3.encoding.homoglyph_greek.trailing.d1; dataset-anger-3.encoding.homoglyph_greek.trailing.d3; dataset-fear-0.encoding.homoglyph_greek.leading.d1; dataset-fear-0.encoding.homoglyph_greek.leading.d3; dataset-fear-0.encoding.homoglyph_greek.middle.d3; dataset-fear-0.encoding.homoglyph_greek.trailing.d3; dataset-fear-3.encoding.homoglyph_greek.leading.d3; dataset-fear-4.encoding.homoglyph_greek.leading.d1; dataset-fear-4.encoding.homoglyph_greek.leading.d3; dataset-joy-0.encoding.homoglyph_greek.leading.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d1; dataset-sadness-0.encoding.homoglyph_greek.middle.d3; dataset-sadness-5.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.leading.d1; handwritten-neutral-2.encoding.homoglyph_greek.leading.d3; handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3; handwritten-neutral-3.encoding.homoglyph_greek.leading.d3; handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3;

perturbation/capitalization: prediction flip

Perturbation / capitalization · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/capitalization produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 7 cases

dataset-anger-1.perturbation.capitalization; dataset-anger-3.perturbation.capitalization; dataset-fear-0.perturbation.capitalization; dataset-fear-3.perturbation.capitalization; dataset-joy-0.perturbation.capitalization; dataset-sadness-3.perturbation.capitalization; dataset-sadness-4.perturbation.capitalization;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.capitalization; dataset-anger-3.perturbation.capitalization; dataset-fear-0.perturbation.capitalization; dataset-fear-3.perturbation.capitalization; dataset-joy-0.perturbation.capitalization; dataset-sadness-3.perturbation.capitalization; dataset-sadness-4.perturbation.capitalization;

F-v2-006 **Medium · 59.8/100** Remaining

perturbation/delete_char: prediction flip

Perturbation / delete_char · prediction flip · 5 affected cases · SILVER

5 case(s) in perturbation/delete_char produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 5 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 5 cases

dataset-anger-1.perturbation.delete_char; dataset-anger-2.perturbation.delete_char; dataset-anger-3.perturbation.delete_char; dataset-fear-0.perturbation.delete_char; dataset-fear-3.perturbation.delete_char;

V1 → V2 evidence · 5 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.perturbation.delete_char; dataset-anger-2.perturbation.delete_char; dataset-anger-3.perturbation.delete_char; dataset-fear-0.perturbation.delete_char; dataset-fear-3.perturbation.delete_char;

F-v2-007 **Medium · 59.8/100** Remaining

perturbation/insert_char: prediction flip

Perturbation / insert_char · prediction flip · 6 affected cases · SILVER

6 case(s) in perturbation/insert_char produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 6 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 6 cases

dataset-anger-1.perturbation.insert_char; dataset-anger-2.perturbation.insert_char; dataset-anger-3.perturbation.insert_char; dataset-fear-0.perturbation.insert_char; dataset-fear-3.perturbation.insert_char; dataset-joy-0.perturbation.insert_char;

V1 → V2 evidence · 6 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 6 cases

dataset-anger-1.perturbation.insert_char; dataset-anger-2.perturbation.insert_char; dataset-anger-3.perturbation.insert_char; dataset-fear-0.perturbation.insert_char; dataset-fear-3.perturbation.insert_char; dataset-joy-0.perturbation.insert_char;

F-v2-008 **Medium · 59.8/100** Remaining

perturbation/keyboard_typo: prediction flip

Perturbation / keyboard_typo · prediction flip · 5 affected cases · SILVER

5 case(s) in perturbation/keyboard_typo produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 5 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 5 cases

dataset-anger-1.perturbation.keyboard_typo; dataset-anger-3.perturbation.keyboard_typo; dataset-fear-0.perturbation.keyboard_typo; dataset-fear-3.perturbation.keyboard_typo; dataset-joy-0.perturbation.keyboard_typo;

V1 → V2 evidence · 5 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.perturbation.keyboard_typo; dataset-anger-3.perturbation.keyboard_typo; dataset-fear-0.perturbation.keyboard_typo; dataset-fear-3.perturbation.keyboard_typo; dataset-joy-0.perturbation.keyboard_typo;

F-v2-011 **Medium · 59.8/100** Remaining

perturbation/swap_adjacent: prediction flip

Perturbation / swap_adjacent · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/swap_adjacent produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 7 cases

dataset-anger-1.perturbation.swap_adjacent; dataset-anger-2.perturbation.swap_adjacent; dataset-anger-3.perturbation.swap_adjacent; dataset-fear-0.perturbation.swap_adjacent; dataset-fear-3.perturbation.swap_adjacent; dataset-joy-0.perturbation.swap_adjacent; dataset-joy-4.perturbation.swap_adjacent;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.swap_adjacent; dataset-anger-2.perturbation.swap_adjacent; dataset-anger-3.perturbation.swap_adjacent; dataset-fear-0.perturbation.swap_adjacent; dataset-fear-3.perturbation.swap_adjacent; dataset-joy-0.perturbation.swap_adjacent; dataset-joy-4.perturbation.swap_adjacent;

F-v2-012 **Medium · 59.8/100** Remaining

perturbation/word_split: prediction flip

Perturbation / word_split · prediction flip · 7 affected cases · SILVER

7 case(s) in perturbation/word_split produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 7 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 7 cases

dataset-anger-1.perturbation.word_split; dataset-anger-2.perturbation.word_split; dataset-anger-3.perturbation.word_split; dataset-fear-0.perturbation.word_split; dataset-fear-3.perturbation.word_split; dataset-joy-0.perturbation.word_split; dataset-joy-4.perturbation.word_split;

V1 → V2 evidence · 7 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 7 cases

dataset-anger-1.perturbation.word_split; dataset-anger-2.perturbation.word_split; dataset-anger-3.perturbation.word_split; dataset-fear-0.perturbation.word_split; dataset-fear-3.perturbation.word_split; dataset-joy-0.perturbation.word_split; dataset-joy-4.perturbation.word_split;

F-v2-009 **Medium · 59.7/100** Remaining

perturbation/keyboard_typo_second_word: prediction flip

Perturbation / keyboard_typo_second_word · prediction flip · 2 affected cases · SILVER

2 case(s) in perturbation/keyboard_typo_second_word produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 2 cases

dataset-sadness-0.perturbation.keyboard_typo_second_word; dataset-sadness-5.perturbation.keyboard_typo_second_word;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-sadness-0.perturbation.keyboard_typo_second_word; dataset-sadness-5.perturbation.keyboard_typo_second_word;

F-v2-001 **Medium · 59.7/100** Remaining

encoding/deletion_backspace: prediction flip

Encoding / deletion_backspace · prediction flip · 5 affected cases · SILVER

5 case(s) in encoding/deletion_backspace produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 5 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 5 cases

dataset-anger-1.encoding.deletion_backspace.leading.d1; dataset-anger-3.encoding.deletion_backspace.trailing.d1; dataset-joy-0.encoding.deletion_backspace.leading.d1; dataset-sadness-5.encoding.deletion_backspace.leading.d1; dataset-surprise-2.encoding.deletion_backspace.trailing.d1;

V1 → V2 evidence · 5 remaining / 2 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 5 cases

dataset-anger-1.encoding.deletion_backspace.leading.d1; dataset-anger-3.encoding.deletion_backspace.trailing.d1; dataset-joy-0.encoding.deletion_backspace.leading.d1; dataset-sadness-5.encoding.deletion_backspace.leading.d1; dataset-surprise-2.encoding.deletion_backspace.trailing.d1;

+ Improved original cases · 2 cases

dataset-sadness-5.encoding.deletion_backspace.trailing.d1; handwritten-neutral-5.encoding.deletion_backspace.middle.d1;

perturbation/misspell: prediction flip

Perturbation / misspell · prediction flip · 6 affected cases · SILVER

6 case(s) in perturbation/misspell produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 6 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 6 cases

dataset-anger-1.perturbation.misspell; dataset-anger-2.perturbation.misspell; dataset-anger-3.perturbation.misspell; dataset-fear-3.perturbation.misspell; dataset-joy-0.perturbation.misspell; dataset-joy-4.perturbation.misspell;

V1 → V2 evidence · 6 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 6 cases

dataset-anger-1.perturbation.misspell; dataset-anger-2.perturbation.misspell; dataset-anger-3.perturbation.misspell; dataset-fear-3.perturbation.misspell; dataset-joy-0.perturbation.misspell; dataset-joy-4.perturbation.misspell;

truncation/signal_end_short: prediction flip

Truncation / signal_end_short · prediction flip · 2 affected cases · SILVER

2 case(s) in truncation/signal_end_short produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 2 cases

dataset-anger-1.truncation.signal_end_short; dataset-sadness-5.truncation.signal_end_short;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-anger-1.truncation.signal_end_short; dataset-sadness-5.truncation.signal_end_short;

F-v2-014 **Medium · 50.5/100** Remaining

truncation/signal_start_short: prediction flip

Truncation / signal_start_short · prediction flip · 2 affected cases · SILVER

2 case(s) in truncation/signal_start_short produced prediction flip -- the top label changed under a meaning-preserving transformation while the clean prediction was confident (at or above 0.6). Severity is scored from the single highest-scoring instance; the 2 affected case(s) indicate reproducibility, not impact.

SILVER qualification. Mechanically justified evidence; semantic equivalence has not been individually hand-verified. Treat meaning-preservation claims with this limitation.

RECOMMENDED REMEDIATION

Model robustness gap rather than an endpoint defect; consider augmenting training data with this perturbation family.

+ Supporting attack IDs · V2 · 2 cases

dataset-anger-1.truncation.signal_start_short; dataset-sadness-5.truncation.signal_start_short;

V1 → V2 evidence · 2 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 2 cases

dataset-anger-1.truncation.signal_start_short; dataset-sadness-5.truncation.signal_start_short;

F-v2-004 **Low · 33.7/100** Remaining

malformed/oversized_10mb: slow response

Malformed / oversized_10mb · slow response · 1 affected case · GOLD

1 case(s) in malformed/oversized_10mb produced slow response -- the response landed in the slow band (2 seconds up to the timeout). Severity is scored from the single highest-scoring instance; the 1 affected case(s) indicate reproducibility, not impact.

RECOMMENDED REMEDIATION

Profile and bound the slow code path so typical requests stay under the 2-second latency band.

+ Supporting attack IDs · V2 · 1 case

malformed.oversized_10mb;

V1 → V2 evidence · 1 remaining / 0 improved / 0 inconclusive. Improved cases concern this failure mode only.

+ Remaining on V2 · 1 case

malformed.oversized_10mb;

Read the result in context.

Clean baseline → Attack library → V1 / V2 runner → Manifest + results → Analysis → This report

Four defenses, one shared model

1. **Length validation.** Reject sequences over the model's limit before inference.
2. **Strict request types.** Reject wrong types and unexpected fields.
3. **Scoped exception handling.** Contain unexpected errors while preserving validation responses.
4. **Unicode and whitespace cleanup.** Normalize text before inference.

No model retraining; no silent truncation. Prediction timeout: 10 seconds. Latency bands: under 2 seconds normal, 2 seconds up to the deadline slow, deadline reached timeout.

Interpretation limits

- Drift is measured against each endpoint's own clean prediction, never the dataset label a baseline sentence was written to represent. Where a baseline does not predict its intended label the comparison remains valid, and no model-accuracy claim in this analysis rests on those baselines.
- V2's defenses are application-layer only; no retraining was done. V2 is therefore not expected to reduce model-level sensitivity to typos or other meaning-preserving perturbations. A perturbation finding that persists on V2 is an honest non-improvement, not a gap in the hardening.
- A failed health check records observed unavailability only. It is not evidence that the operating-system process died, and no process-death claim is made anywhere here.
- Validity tiers are preserved per finding. Most of the suite is SILVER -- mechanically justified but not hand-verified -- so findings resting on SILVER evidence carry weaker semantic assurance than GOLD ones.
- Category failure rates and prediction-flip rates use different eligible denominators. A zero denominator is displayed as N/A.
- Manifest and suite compatibility are supplied by analysis. The report validates the JSON interface; it does not independently replay requests or prove when the manifest was registered.
- Timeouts establish that the request deadline was reached. These artifacts do not identify which server processing stage caused the delay.
- Leak detection is based on signatures and does not establish the absence of all possible information leaks.

Trace every qualification.

Expand sections in the browser to inspect case-level evidence. The PDF prints the complete appendix.

+ V1 denominators, exclusions and confidence evidence

Category eligibility

CATEGORY	FAILED / ELIGIBLE	RATE	REVIEW	DIAGNOSTIC	UNEVALUABLE
Malformed	4 / 21	19.05%	0	0	0
Boundary	1 / 11	9.09%	0	0	0
Perturbation	45 / 377	11.94%	0	0	0
Encoding	104 / 931	11.17%	42	0	0
Whitespace	68 / 210	32.38%	0	42	0
Truncation	88 / 168	52.38%	0	84	0

Baselines below the confidence threshold

BASELINE ID	CLEAN CONFIDENCE
handwritten-neutral-4	0.4655
handwritten-neutral-0	0.4688
dataset-surprise-0	0.4906
handwritten-disgust-4	0.5072
handwritten-disgust-2	0.5228
dataset-surprise-5	0.5481

Drift exclusions and unavailable evidence

Excluded · review unresolved	42
Excluded · oracle not label match baseline	210
Excluded · low clean confidence	228
Unevaluable · missing attacked prediction	3

+ Per-flip confidence evidence · 218 cases

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-anger-0.encoding.compatibility_fullwidth.whole	0.9910	0.4389	-0.5521
dataset-anger-0.whitespace.interior_runs	0.9910	0.5716	-0.4194
dataset-anger-1.perturbation.swap_adjacent	0.6179	0.9431	+0.3252
dataset-anger-1.perturbation.keyboard_type	0.6179	0.8352	+0.2173
dataset-anger-1.perturbation.delete_char	0.6179	0.8551	+0.2372
dataset-anger-1.perturbation.word_split	0.6179	0.8813	+0.2634
dataset-anger-1.perturbation.insert_char	0.6179	0.9664	+0.3485
dataset-anger-1.perturbation.misspell	0.6179	0.8283	+0.2104
dataset-anger-1.perturbation.capitalization	0.6179	0.9740	+0.3561
dataset-anger-1.encoding.homoglyph.leading.d3	0.6179	0.7187	+0.1008

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-anger-1.encoding.homoglyph_greek.leading.d1	0.6179	0.8425	+0.2246
dataset-anger-1.encoding.homoglyph_greek.leading.d3	0.6179	0.6035	-0.0144
dataset-anger-1.encoding.invisible_zero_width.leading.d3	0.6179	0.7557	+0.1378
dataset-anger-1.encoding.deletion_backspace.leading.d1	0.6179	0.8368	+0.2189
dataset-anger-1.encoding.compatibility_fullwidth.whole	0.6179	0.4567	-0.1612
dataset-anger-1.whitespace.interior_runs	0.6179	0.6004	-0.0175
dataset-anger-1.whitespace.tabs	0.6179	0.9695	+0.3516
dataset-anger-1.whitespace.newlines	0.6179	0.9925	+0.3746
dataset-anger-1.whitespace.unicode_spaces	0.6179	0.5117	-0.1062
dataset-anger-1.truncation.signal_start_short	0.6179	0.5982	-0.0197
dataset-anger-1.truncation.signal_end_short	0.6179	0.8986	+0.2807
dataset-anger-2.perturbation.swap_adjacent	0.9792	0.8506	-0.1286
dataset-anger-2.perturbation.delete_char	0.9792	0.4867	-0.4925
dataset-anger-2.perturbation.word_split	0.9792	0.4736	-0.5055
dataset-anger-2.perturbation.insert_char	0.9792	0.8988	-0.0804
dataset-anger-2.perturbation.misspell	0.9792	0.9339	-0.0453
dataset-anger-2.encoding.homoglyph.leading.d3	0.9792	0.5646	-0.4146
dataset-anger-2.encoding.homoglyph.trailing.d3	0.9792	0.9258	-0.0534
dataset-anger-2.encoding.homoglyph.trailing.d5	0.9792	0.9600	-0.0191
dataset-anger-2.encoding.compatibility_fullwidth.whole	0.9792	0.4806	-0.4985
dataset-anger-2.whitespace.interior_runs	0.9792	0.4263	-0.5529
dataset-anger-2.whitespace.tabs	0.9792	0.9383	-0.0409
dataset-anger-2.whitespace.newlines	0.9792	0.9314	-0.0478
dataset-anger-2.whitespace.unicode_spaces	0.9792	0.8682	-0.1110
dataset-anger-3.perturbation.swap_adjacent	0.9945	0.8996	-0.0949
dataset-anger-3.perturbation.keyboard_typo	0.9945	0.9822	-0.0123
dataset-anger-3.perturbation.delete_char	0.9945	0.9612	-0.0334
dataset-anger-3.perturbation.word_split	0.9945	0.7544	-0.2401
dataset-anger-3.perturbation.insert_char	0.9945	0.9081	-0.0864
dataset-anger-3.perturbation.misspell	0.9945	0.5463	-0.4482
dataset-anger-3.perturbation.capitalization	0.9945	0.7199	-0.2746
dataset-anger-3.encoding.homoglyph.trailing.d1	0.9945	0.9004	-0.0941
dataset-anger-3.encoding.homoglyph.trailing.d3	0.9945	0.4974	-0.4971
dataset-anger-3.encoding.homoglyph.trailing.d5	0.9945	0.5557	-0.4388
dataset-anger-3.encoding.homoglyph_greek.middle.d3	0.9945	0.7024	-0.2921
dataset-anger-3.encoding.homoglyph_greek.trailing.d1	0.9945	0.6885	-0.3060
dataset-anger-3.encoding.homoglyph_greek.trailing.d3	0.9945	0.7024	-0.2921
dataset-anger-3.encoding.invisible_zero_width.trailing.d3	0.9945	0.7031	-0.2914
dataset-anger-3.encoding.deletion_backspace.trailing.d1	0.9945	0.9769	-0.0177
dataset-anger-3.encoding.compatibility_fullwidth.whole	0.9945	0.5432	-0.4513
dataset-anger-4.encoding.compatibility_fullwidth.whole	0.9954	0.4613	-0.5341
dataset-anger-4.whitespace.interior_runs	0.9954	0.4717	-0.5236
dataset-anger-5.whitespace.interior_runs	0.9752	0.5894	-0.3859

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-anger-5.whitespace.tabs	0.9752	0.4211	-0.5542
dataset-anger-5.whitespace.newlines	0.9752	0.8124	-0.1629
dataset-fear-0.perturbation.swap_adjacent	0.9960	0.7988	-0.1972
dataset-fear-0.perturbation.keyboard_typo	0.9960	0.5984	-0.3976
dataset-fear-0.perturbation.delete_char	0.9960	0.8895	-0.1065
dataset-fear-0.perturbation.word_split	0.9960	0.7092	-0.2868
dataset-fear-0.perturbation.insert_char	0.9960	0.9559	-0.0401
dataset-fear-0.perturbation.capitalization	0.9960	0.8950	-0.1010
dataset-fear-0.encoding.homoglyph.leading.d5	0.9960	0.6777	-0.3183
dataset-fear-0.encoding.homoglyph.middle.d1	0.9960	0.8209	-0.1751
dataset-fear-0.encoding.homoglyph.middle.d3	0.9960	0.9415	-0.0545
dataset-fear-0.encoding.homoglyph.middle.d5	0.9960	0.7840	-0.2120
dataset-fear-0.encoding.homoglyph.trailing.d5	0.9960	0.9467	-0.0492
dataset-fear-0.encoding.homoglyph_greek.leading.d1	0.9960	0.7171	-0.2789
dataset-fear-0.encoding.homoglyph_greek.leading.d3	0.9960	0.7314	-0.2645
dataset-fear-0.encoding.homoglyph_greek.middle.d3	0.9960	0.7314	-0.2645
dataset-fear-0.encoding.homoglyph_greek.trailing.d3	0.9960	0.7314	-0.2645
dataset-fear-0.encoding.compatibility_fullwidth.whole	0.9960	0.5556	-0.4404
dataset-fear-0.whitespace.tabs	0.9960	0.5296	-0.4664
dataset-fear-1.encoding.compatibility_fullwidth.whole	0.9310	0.4327	-0.4983
dataset-fear-1.whitespace.tabs	0.9310	0.6498	-0.2811
dataset-fear-1.whitespace.newlines	0.9310	0.6961	-0.2349
dataset-fear-1.whitespace.leading_trailing	0.9310	0.5585	-0.3725
dataset-fear-1.whitespace.unicode_spaces	0.9310	0.7704	-0.1606
dataset-fear-2.encoding.compatibility_fullwidth.whole	0.9957	0.4325	-0.5632
dataset-fear-3.perturbation.swap_adjacent	0.9670	0.9935	+0.0265
dataset-fear-3.perturbation.keyboard_typo	0.9670	0.9940	+0.0270
dataset-fear-3.perturbation.delete_char	0.9670	0.9946	+0.0276
dataset-fear-3.perturbation.word_split	0.9670	0.9944	+0.0274
dataset-fear-3.perturbation.insert_char	0.9670	0.9939	+0.0269
dataset-fear-3.perturbation.misspell	0.9670	0.9938	+0.0267
dataset-fear-3.perturbation.capitalization	0.9670	0.9939	+0.0268
dataset-fear-3.encoding.homoglyph_greek.leading.d3	0.9670	0.9944	+0.0274
dataset-fear-3.encoding.compatibility_fullwidth.whole	0.9670	0.4606	-0.5065
dataset-fear-3.whitespace.interior_runs	0.9670	0.3291	-0.6379
dataset-fear-3.whitespace.tabs	0.9670	0.3935	-0.5735
dataset-fear-3.whitespace.newlines	0.9670	0.8781	-0.0890
dataset-fear-3.whitespace.unicode_spaces	0.9670	0.9389	-0.0281
dataset-fear-4.encoding.homoglyph.leading.d5	0.9947	0.7125	-0.2822
dataset-fear-4.encoding.homoglyph_greek.leading.d1	0.9947	0.8606	-0.1341
dataset-fear-4.encoding.homoglyph_greek.leading.d3	0.9947	0.9900	-0.0046
dataset-fear-4.encoding.compatibility_fullwidth.whole	0.9947	0.4596	-0.5351
dataset-fear-4.whitespace.interior_runs	0.9947	0.4080	-0.5867

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-fear-5.encoding.compatibility_fullwidth.whole	0.9954	0.4574	-0.5380
dataset-joy-0.perturbation.swap_adjacent	0.8612	0.4657	-0.3955
dataset-joy-0.perturbation.keyboard_typo	0.8612	0.5581	-0.3031
dataset-joy-0.perturbation.word_split	0.8612	0.5758	-0.2854
dataset-joy-0.perturbation.insert_char	0.8612	0.5145	-0.3467
dataset-joy-0.perturbation.misspell	0.8612	0.5171	-0.3441
dataset-joy-0.perturbation.capitalization	0.8612	0.5332	-0.3279
dataset-joy-0.encoding.homoglyph.leading.d1	0.8612	0.4867	-0.3745
dataset-joy-0.encoding.homoglyph.leading.d3	0.8612	0.7629	-0.0983
dataset-joy-0.encoding.homoglyph.leading.d5	0.8612	0.4791	-0.3820
dataset-joy-0.encoding.homoglyph_greek.leading.d1	0.8612	0.5158	-0.3454
dataset-joy-0.encoding.deletion_backspace.leading.d1	0.8612	0.7507	-0.1105
dataset-joy-0.encoding.compatibility_fullwidth.whole	0.8612	0.4386	-0.4226
dataset-joy-0.whitespace.leading_trailing	0.8612	0.4795	-0.3817
dataset-joy-1.whitespace.tabs	0.9942	0.8058	-0.1884
dataset-joy-1.whitespace.newlines	0.9942	0.7417	-0.2525
dataset-joy-1.whitespace.unicode_spaces	0.9942	0.6130	-0.3811
dataset-joy-2.encoding.compatibility_fullwidth.whole	0.9937	0.4290	-0.5647
dataset-joy-3.encoding.homoglyph.trailing.d3	0.9882	0.9630	-0.0253
dataset-joy-3.encoding.homoglyph.trailing.d5	0.9882	0.8401	-0.1481
dataset-joy-3.encoding.compatibility_fullwidth.whole	0.9882	0.4357	-0.5525
dataset-joy-3.whitespace.tabs	0.9882	0.6420	-0.3462
dataset-joy-3.whitespace.newlines	0.9882	0.6546	-0.3336
dataset-joy-4.perturbation.swap_adjacent	0.9919	0.9588	-0.0331
dataset-joy-4.perturbation.word_split	0.9919	0.5962	-0.3958
dataset-joy-4.perturbation.misspell	0.9919	0.6670	-0.3249
dataset-joy-4.encoding.compatibility_fullwidth.whole	0.9919	0.3995	-0.5925
dataset-joy-5.encoding.homoglyph.trailing.d5	0.9927	0.8546	-0.1382
dataset-joy-5.encoding.compatibility_fullwidth.whole	0.9927	0.5602	-0.4326
dataset-joy-5.whitespace.tabs	0.9927	0.4063	-0.5864
dataset-sadness-0.perturbation.keyboard_typo_second_word	0.9946	0.8639	-0.1307
dataset-sadness-0.encoding.homoglyph.middle.d5	0.9946	0.6602	-0.3344
dataset-sadness-0.encoding.homoglyph_greek.middle.d1	0.9946	0.9840	-0.0106
dataset-sadness-0.encoding.homoglyph_greek.middle.d3	0.9946	0.9841	-0.0105
dataset-sadness-0.encoding.compatibility_fullwidth.whole	0.9946	0.4755	-0.5191
dataset-sadness-0.whitespace.interior_runs	0.9946	0.4238	-0.5708
dataset-sadness-1.encoding.compatibility_fullwidth.whole	0.9935	0.4783	-0.5152
dataset-sadness-2.encoding.compatibility_fullwidth.whole	0.9920	0.4419	-0.5501
dataset-sadness-3.perturbation.capitalization	0.9940	0.9849	-0.0091
dataset-sadness-3.encoding.homoglyph.leading.d5	0.9940	0.2818	-0.7122
dataset-sadness-3.encoding.homoglyph.middle.d5	0.9940	0.5240	-0.4700
dataset-sadness-3.encoding.homoglyph.trailing.d5	0.9940	0.9343	-0.0597
dataset-sadness-3.encoding.compatibility_fullwidth.whole	0.9940	0.5964	-0.3976

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-sadness-4.perturbation.capitalization	0.9925	0.9916	-0.0009
dataset-sadness-4.encoding.homoglyph.leading.d5	0.9925	0.7658	-0.2267
dataset-sadness-4.encoding.compatibility_fullwidth.whole	0.9925	0.5147	-0.4778
dataset-sadness-5.perturbation.keyboard_typo_second_word	0.8097	0.6541	-0.1556
dataset-sadness-5.encoding.homoglyph.leading.d3	0.8097	0.7401	-0.0696
dataset-sadness-5.encoding.homoglyph.leading.d5	0.8097	0.8031	-0.0066
dataset-sadness-5.encoding.homoglyph_greek.leading.d3	0.8097	0.7544	-0.0553
dataset-sadness-5.encoding.invisible_zero_width.leading.d3	0.8097	0.5791	-0.2306
dataset-sadness-5.encoding.deletion_backspace.leading.d1	0.8097	0.7374	-0.0723
dataset-sadness-5.encoding.deletion_backspace.trailing.d1	0.8097	0.5695	-0.2402
dataset-sadness-5.encoding.compatibility_fullwidth.whole	0.8097	0.4149	-0.3948
dataset-sadness-5.whitespace.tabs	0.8097	0.8920	+0.0823
dataset-sadness-5.whitespace.newlines	0.8097	0.7807	-0.0289
dataset-sadness-5.whitespace.leading_trailing	0.8097	0.5324	-0.2773
dataset-sadness-5.truncation.signal_start_short	0.8097	0.6834	-0.1263
dataset-sadness-5.truncation.signal_end_short	0.8097	0.6836	-0.1261
dataset-surprise-1.encoding.compatibility_fullwidth.whole	0.9837	0.4318	-0.5519
dataset-surprise-1.whitespace.tabs	0.9837	0.5135	-0.4702
dataset-surprise-1.whitespace.newlines	0.9837	0.7382	-0.2455
dataset-surprise-1.whitespace.unicode_spaces	0.9837	0.4853	-0.4984
dataset-surprise-2.encoding.homoglyph.trailing.d1	0.9773	0.9729	-0.0044
dataset-surprise-2.encoding.homoglyph.trailing.d3	0.9773	0.9942	+0.0169
dataset-surprise-2.encoding.homoglyph.trailing.d5	0.9773	0.9933	+0.0160
dataset-surprise-2.encoding.invisible_zero_width.trailing.d3	0.9773	0.9928	+0.0155
dataset-surprise-2.encoding.deletion_backspace.trailing.d1	0.9773	0.9912	+0.0139
dataset-surprise-2.encoding.compatibility_fullwidth.whole	0.9773	0.5367	-0.4406
dataset-surprise-2.whitespace.interior_runs	0.9773	0.7356	-0.2417
dataset-surprise-2.whitespace.tabs	0.9773	0.9863	+0.0089
dataset-surprise-2.whitespace.newlines	0.9773	0.9925	+0.0152
dataset-surprise-2.whitespace.unicode_spaces	0.9773	0.9924	+0.0150
dataset-surprise-3.whitespace.interior_runs	0.8638	0.5761	-0.2877
dataset-surprise-3.whitespace.newlines	0.8638	0.7061	-0.1577
dataset-surprise-3.whitespace.leading_trailing	0.8638	0.5531	-0.3107
dataset-surprise-4.encoding.compatibility_fullwidth.whole	0.9863	0.4299	-0.5564
dataset-surprise-4.whitespace.interior_runs	0.9863	0.4030	-0.5833
handwritten-disgust-0.encoding.compatibility_fullwidth.whole	0.9762	0.4287	-0.5474
handwritten-disgust-0.whitespace.interior_runs	0.9762	0.6086	-0.3676
handwritten-disgust-0.whitespace.tabs	0.9762	0.3858	-0.5904
handwritten-disgust-0.whitespace.unicode_spaces	0.9762	0.3080	-0.6682
handwritten-disgust-1.encoding.compatibility_fullwidth.whole	0.9866	0.4713	-0.5153
handwritten-disgust-1.whitespace.interior_runs	0.9866	0.4024	-0.5842
handwritten-disgust-1.whitespace.tabs	0.9866	0.3891	-0.5975
handwritten-disgust-1.whitespace.unicode_spaces	0.9866	0.4998	-0.4868

ATTACK ID	CLEAN	ATTACKED	DELTA
handwritten-disgust-3.encoding.compatibility_fullwidth.whole	0.9843	0.4629	-0.5214
handwritten-disgust-3.whitespace.interior_runs	0.9843	0.4716	-0.5127
handwritten-disgust-3.whitespace.tabs	0.9843	0.3123	-0.6720
handwritten-disgust-5.encoding.compatibility_fullwidth.whole	0.9890	0.4265	-0.5626
handwritten-disgust-5.whitespace.interior_runs	0.9890	0.5255	-0.4636
handwritten-disgust-5.whitespace.tabs	0.9890	0.7800	-0.2090
handwritten-disgust-5.whitespace.unicode_spaces	0.9890	0.6691	-0.3199
handwritten-neutral-1.encoding.homoglyph.middle.d3	0.7554	0.2680	-0.4874
handwritten-neutral-1.encoding.homoglyph.middle.d5	0.7554	0.4361	-0.3192
handwritten-neutral-1.whitespace.interior_runs	0.7554	0.3748	-0.3806
handwritten-neutral-1.whitespace.leading_trailing	0.7554	0.5479	-0.2074
handwritten-neutral-2.encoding.homoglyph.middle.d5	0.8379	0.4958	-0.3421
handwritten-neutral-2.encoding.homoglyph_greek.leading.d1	0.8379	0.7181	-0.1198
handwritten-neutral-2.encoding.homoglyph_greek.leading.d3	0.8379	0.6184	-0.2195
handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3	0.8379	0.4022	-0.4357
handwritten-neutral-2.encoding.invisible_zero_width.leading.d3	0.8379	0.6287	-0.2092
handwritten-neutral-2.whitespace.interior_runs	0.8379	0.3513	-0.4866
handwritten-neutral-2.whitespace.leading_trailing	0.8379	0.5395	-0.2984
handwritten-neutral-3.encoding.homoglyph_greek.leading.d3	0.7580	0.3647	-0.3932
handwritten-neutral-3.whitespace.interior_runs	0.7580	0.4219	-0.3361
handwritten-neutral-3.whitespace.tabs	0.7580	0.4712	-0.2868
handwritten-neutral-3.whitespace.newlines	0.7580	0.3817	-0.3763
handwritten-neutral-3.whitespace.leading_trailing	0.7580	0.5088	-0.2491
handwritten-neutral-5.encoding.homoglyph.middle.d1	0.7872	0.5077	-0.2795
handwritten-neutral-5.encoding.homoglyph.middle.d3	0.7872	0.3702	-0.4170
handwritten-neutral-5.encoding.homoglyph.middle.d5	0.7872	0.6500	-0.1372
handwritten-neutral-5.encoding.homoglyph.trailing.d1	0.7872	0.5308	-0.2564
handwritten-neutral-5.encoding.homoglyph.trailing.d5	0.7872	0.4990	-0.2882
handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3	0.7872	0.9451	+0.1579
handwritten-neutral-5.encoding.invisible_zero_width.middle.d3	0.7872	0.9082	+0.1210
handwritten-neutral-5.encoding.invisible_zero_width.trailing.d3	0.7872	0.4674	-0.3197
handwritten-neutral-5.encoding.deletion_backspace.middle.d1	0.7872	0.5173	-0.2698
handwritten-neutral-5.whitespace.interior_runs	0.7872	0.3422	-0.4449
handwritten-neutral-5.whitespace.tabs	0.7872	0.5489	-0.2383
handwritten-neutral-5.whitespace.newlines	0.7872	0.3763	-0.4109
handwritten-neutral-5.whitespace.leading_trailing	0.7872	0.4632	-0.3240

+ V1 evidence validity and unscored observations

Finding groups by validity tier: GOLD 8 · SILVER 21. Mixed-tier groups contribute to each applicable tier; these are group counts.

Review queue · 42 cases

Unresolved REVIEW cases are excluded from automatic vulnerability scoring.

dataset-anger-0.encoding.bidi_reorder.middle; dataset-anger-1.encoding.bidi_reorder.middle; dataset-anger-2.encoding.bidi_reorder.middle; dataset-anger-3.encoding.bidi_reorder.middle; dataset-anger-4.encoding.bidi_reorder.middle; dataset-anger-5.encoding.bidi_reorder.middle; dataset-fear-0.encoding.bidi_reorder.middle; dataset-fear-1.encoding.bidi_reorder.middle; dataset-fear-2.encoding.bidi_reorder.middle; dataset-fear-3.encoding.bidi_reorder.middle; dataset-fear-4.encoding.bidi_reorder.middle; dataset-fear-5.encoding.bidi_reorder.middle; dataset-joy-0.encoding.bidi_reorder.middle; dataset-joy-1.encoding.bidi_reorder.middle; dataset-joy-2.encoding.bidi_reorder.middle; dataset-joy-3.encoding.bidi_reorder.middle; dataset-joy-4.encoding.bidi_reorder.middle; dataset-joy-5.encoding.bidi_reorder.middle; dataset-sadness-0.encoding.bidi_reorder.middle; dataset-sadness-1.encoding.bidi_reorder.middle; dataset-sadness-2.encoding.bidi_reorder.middle; dataset-sadness-3.encoding.bidi_reorder.middle; dataset-sadness-4.encoding.bidi_reorder.middle; dataset-sadness-5.encoding.bidi_reorder.middle; dataset-surprise-0.encoding.bidi_reorder.middle; dataset-surprise-1.encoding.bidi_reorder.middle; dataset-surprise-2.encoding.bidi_reorder.middle; dataset-surprise-3.encoding.bidi_reorder.middle; dataset-surprise-4.encoding.bidi_reorder.middle; dataset-surprise-5.encoding.bidi_reorder.middle; handwritten-disgust-0.encoding.bidi_reorder.middle; handwritten-disgust-1.encoding.bidi_reorder.middle; handwritten-disgust-2.encoding.bidi_reorder.middle; handwritten-disgust-3.encoding.bidi_reorder.middle; handwritten-disgust-4.encoding.bidi_reorder.middle; handwritten-disgust-5.encoding.bidi_reorder.middle; handwritten-neutral-0.encoding.bidi_reorder.middle; handwritten-neutral-1.encoding.bidi_reorder.middle; handwritten-neutral-2.encoding.bidi_reorder.middle; handwritten-neutral-3.encoding.bidi_reorder.middle; handwritten-neutral-4.encoding.bidi_reorder.middle; handwritten-neutral-5.encoding.bidi_reorder.middle;

Diagnostic observations · 126 cases

Diagnostic oracles are unscored regardless of their recognized validity tier. Tier differences are preserved below.

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-anger-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-3.truncation.signal_tail_only	DIAGNOSTIC	No

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-fear-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-0.whitespace.repeated_punctuation	REVIEW	Yes

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-surprise-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-5.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-0.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-0.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-0.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-1.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-1.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-1.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-2.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-2.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-2.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-3.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-3.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-3.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-4.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-4.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-4.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-5.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-5.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-5.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-0.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-0.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-0.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-1.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-1.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-1.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-2.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-2.truncation.signal_head_only	DIAGNOSTIC	No

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
handwritten-neutral-2.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-3.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-3.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-3.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-4.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-4.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-4.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-5.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-5.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-5.truncation.signal_tail_only	DIAGNOSTIC	No

+ V2 denominators, exclusions and confidence evidence

Category eligibility

CATEGORY	FAILED / ELIGIBLE	RATE	REVIEW	DIAGNOSTIC	UNEVALUABLE
Malformed	1 / 21	4.76%	0	0	0
Boundary	0 / 11	0.00%	0	0	0
Perturbation	45 / 377	11.94%	0	0	0
Encoding	63 / 931	6.77%	42	0	0
Whitespace	0 / 210	0.00%	0	42	0
Truncation	4 / 168	2.38%	0	84	0

Baselines below the confidence threshold

BASELINE ID	CLEAN CONFIDENCE
handwritten-neutral-4	0.4655
handwritten-neutral-0	0.4688
dataset-surprise-0	0.4906
handwritten-disgust-4	0.5072
handwritten-disgust-2	0.5228
dataset-surprise-5	0.5481

Drift exclusions and unavailable evidence

Excluded · review unresolved	42
Excluded · oracle not label match baseline	210
Excluded · low clean confidence	228

+ Per-flip confidence evidence · 112 cases

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-anger-1.perturbation.swap_adjacent	0.6179	0.9431	+0.3252
dataset-anger-1.perturbation.keyboard_typo	0.6179	0.8352	+0.2173
dataset-anger-1.perturbation.delete_char	0.6179	0.8551	+0.2372
dataset-anger-1.perturbation.word_split	0.6179	0.8813	+0.2634
dataset-anger-1.perturbation.insert_char	0.6179	0.9664	+0.3485

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-anger-1.perturbation.misspell	0.6179	0.8283	+0.2104
dataset-anger-1.perturbation.capitalization	0.6179	0.9740	+0.3561
dataset-anger-1.encoding.homoglyph.leading.d3	0.6179	0.7187	+0.1008
dataset-anger-1.encoding.homoglyph_greek.leading.d1	0.6179	0.8425	+0.2246
dataset-anger-1.encoding.homoglyph_greek.leading.d3	0.6179	0.6035	-0.0144
dataset-anger-1.encoding.deletion_backspace.leading.d1	0.6179	0.8903	+0.2724
dataset-anger-1.truncation.signal_start_short	0.6179	0.5982	-0.0197
dataset-anger-1.truncation.signal_end_short	0.6179	0.8986	+0.2807
dataset-anger-2.perturbation.swap_adjacent	0.9792	0.8506	-0.1286
dataset-anger-2.perturbation.delete_char	0.9792	0.4867	-0.4925
dataset-anger-2.perturbation.word_split	0.9792	0.4736	-0.5055
dataset-anger-2.perturbation.insert_char	0.9792	0.8988	-0.0804
dataset-anger-2.perturbation.misspell	0.9792	0.9339	-0.0453
dataset-anger-2.encoding.homoglyph.leading.d3	0.9792	0.5646	-0.4146
dataset-anger-2.encoding.homoglyph.trailing.d3	0.9792	0.9258	-0.0534
dataset-anger-2.encoding.homoglyph.trailing.d5	0.9792	0.9600	-0.0191
dataset-anger-3.perturbation.swap_adjacent	0.9945	0.8996	-0.0949
dataset-anger-3.perturbation.keyboard_typo	0.9945	0.9822	-0.0123
dataset-anger-3.perturbation.delete_char	0.9945	0.9612	-0.0334
dataset-anger-3.perturbation.word_split	0.9945	0.7544	-0.2401
dataset-anger-3.perturbation.insert_char	0.9945	0.9081	-0.0864
dataset-anger-3.perturbation.misspell	0.9945	0.5463	-0.4482
dataset-anger-3.perturbation.capitalization	0.9945	0.7199	-0.2746
dataset-anger-3.encoding.homoglyph.trailing.d1	0.9945	0.9004	-0.0941
dataset-anger-3.encoding.homoglyph.trailing.d3	0.9945	0.4974	-0.4971
dataset-anger-3.encoding.homoglyph.trailing.d5	0.9945	0.5557	-0.4388
dataset-anger-3.encoding.homoglyph_greek.middle.d3	0.9945	0.7024	-0.2921
dataset-anger-3.encoding.homoglyph_greek.trailing.d1	0.9945	0.6885	-0.3060
dataset-anger-3.encoding.homoglyph_greek.trailing.d3	0.9945	0.7024	-0.2921
dataset-anger-3.encoding.deletion_backspace.trailing.d1	0.9945	0.9249	-0.0696
dataset-fear-0.perturbation.swap_adjacent	0.9960	0.7988	-0.1972
dataset-fear-0.perturbation.keyboard_typo	0.9960	0.5984	-0.3976
dataset-fear-0.perturbation.delete_char	0.9960	0.8895	-0.1065
dataset-fear-0.perturbation.word_split	0.9960	0.7092	-0.2868
dataset-fear-0.perturbation.insert_char	0.9960	0.9559	-0.0401
dataset-fear-0.perturbation.capitalization	0.9960	0.8950	-0.1010
dataset-fear-0.encoding.homoglyph.leading.d5	0.9960	0.6777	-0.3183
dataset-fear-0.encoding.homoglyph.middle.d1	0.9960	0.8209	-0.1751
dataset-fear-0.encoding.homoglyph.middle.d3	0.9960	0.9415	-0.0545
dataset-fear-0.encoding.homoglyph.middle.d5	0.9960	0.7840	-0.2120
dataset-fear-0.encoding.homoglyph.trailing.d5	0.9960	0.9467	-0.0492
dataset-fear-0.encoding.homoglyph_greek.leading.d1	0.9960	0.7171	-0.2789
dataset-fear-0.encoding.homoglyph_greek.leading.d3	0.9960	0.7314	-0.2645

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-fear-0.encoding.homoglyph_greek.middle.d3	0.9960	0.7314	-0.2645
dataset-fear-0.encoding.homoglyph_greek.trailing.d3	0.9960	0.7314	-0.2645
dataset-fear-3.perturbation.swap_adjacent	0.9670	0.9935	+0.0265
dataset-fear-3.perturbation.keyboard_typo	0.9670	0.9940	+0.0270
dataset-fear-3.perturbation.delete_char	0.9670	0.9946	+0.0276
dataset-fear-3.perturbation.word_split	0.9670	0.9944	+0.0274
dataset-fear-3.perturbation.insert_char	0.9670	0.9939	+0.0269
dataset-fear-3.perturbation.misspell	0.9670	0.9938	+0.0267
dataset-fear-3.perturbation.capitalization	0.9670	0.9939	+0.0268
dataset-fear-3.encoding.homoglyph_greek.leading.d3	0.9670	0.9944	+0.0274
dataset-fear-4.encoding.homoglyph.leading.d5	0.9947	0.7125	-0.2822
dataset-fear-4.encoding.homoglyph_greek.leading.d1	0.9947	0.8606	-0.1341
dataset-fear-4.encoding.homoglyph_greek.leading.d3	0.9947	0.9900	-0.0046
dataset-joy-0.perturbation.swap_adjacent	0.8612	0.4657	-0.3955
dataset-joy-0.perturbation.keyboard_typo	0.8612	0.5581	-0.3031
dataset-joy-0.perturbation.word_split	0.8612	0.5758	-0.2854
dataset-joy-0.perturbation.insert_char	0.8612	0.5145	-0.3467
dataset-joy-0.perturbation.misspell	0.8612	0.5171	-0.3441
dataset-joy-0.perturbation.capitalization	0.8612	0.5332	-0.3279
dataset-joy-0.encoding.homoglyph.leading.d1	0.8612	0.4867	-0.3745
dataset-joy-0.encoding.homoglyph.leading.d3	0.8612	0.7629	-0.0983
dataset-joy-0.encoding.homoglyph.leading.d5	0.8612	0.4791	-0.3820
dataset-joy-0.encoding.homoglyph_greek.leading.d1	0.8612	0.5158	-0.3454
dataset-joy-0.encoding.deletion_backspace.leading.d1	0.8612	0.6772	-0.1840
dataset-joy-3.encoding.homoglyph.trailing.d3	0.9882	0.9630	-0.0253
dataset-joy-3.encoding.homoglyph.trailing.d5	0.9882	0.8401	-0.1481
dataset-joy-4.perturbation.swap_adjacent	0.9919	0.9588	-0.0331
dataset-joy-4.perturbation.word_split	0.9919	0.5962	-0.3958
dataset-joy-4.perturbation.misspell	0.9919	0.6670	-0.3249
dataset-joy-5.encoding.homoglyph.trailing.d5	0.9927	0.8546	-0.1382
dataset-sadness-0.perturbation.keyboard_typo_second_word	0.9946	0.8639	-0.1307
dataset-sadness-0.encoding.homoglyph.middle.d5	0.9946	0.6602	-0.3344
dataset-sadness-0.encoding.homoglyph_greek.middle.d1	0.9946	0.9840	-0.0106
dataset-sadness-0.encoding.homoglyph_greek.middle.d3	0.9946	0.9841	-0.0105
dataset-sadness-3.perturbation.capitalization	0.9940	0.9849	-0.0091
dataset-sadness-3.encoding.homoglyph.leading.d5	0.9940	0.2818	-0.7122
dataset-sadness-3.encoding.homoglyph.middle.d5	0.9940	0.5240	-0.4700
dataset-sadness-3.encoding.homoglyph.trailing.d5	0.9940	0.9343	-0.0597
dataset-sadness-4.perturbation.capitalization	0.9925	0.9916	-0.0009
dataset-sadness-4.encoding.homoglyph.leading.d5	0.9925	0.7658	-0.2267
dataset-sadness-5.perturbation.keyboard_typo_second_word	0.8097	0.6541	-0.1556
dataset-sadness-5.encoding.homoglyph.leading.d3	0.8097	0.7401	-0.0696
dataset-sadness-5.encoding.homoglyph.leading.d5	0.8097	0.8031	-0.0066

ATTACK ID	CLEAN	ATTACKED	DELTA
dataset-sadness-5.encoding.homoglyph_greek.leading.d3	0.8097	0.7544	-0.0553
dataset-sadness-5.encoding.deletion_backspace.leading.d1	0.8097	0.5000	-0.3097
dataset-sadness-5.truncation.signal_start_short	0.8097	0.6834	-0.1263
dataset-sadness-5.truncation.signal_end_short	0.8097	0.6836	-0.1261
dataset-surprise-2.encoding.homoglyph.trailing.d1	0.9773	0.9729	-0.0044
dataset-surprise-2.encoding.homoglyph.trailing.d3	0.9773	0.9942	+0.0169
dataset-surprise-2.encoding.homoglyph.trailing.d5	0.9773	0.9933	+0.0160
dataset-surprise-2.encoding.deletion_backspace.trailing.d1	0.9773	0.9733	-0.0040
handwritten-neutral-1.encoding.homoglyph.middle.d3	0.7554	0.2680	-0.4874
handwritten-neutral-1.encoding.homoglyph.middle.d5	0.7554	0.4361	-0.3192
handwritten-neutral-2.encoding.homoglyph.middle.d5	0.8379	0.4958	-0.3421
handwritten-neutral-2.encoding.homoglyph_greek.leading.d1	0.8379	0.7181	-0.1198
handwritten-neutral-2.encoding.homoglyph_greek.leading.d3	0.8379	0.6184	-0.2195
handwritten-neutral-2.encoding.homoglyph_greek.trailing.d3	0.8379	0.4022	-0.4357
handwritten-neutral-3.encoding.homoglyph_greek.leading.d3	0.7580	0.3647	-0.3932
handwritten-neutral-5.encoding.homoglyph.middle.d1	0.7872	0.5077	-0.2795
handwritten-neutral-5.encoding.homoglyph.middle.d3	0.7872	0.3702	-0.4170
handwritten-neutral-5.encoding.homoglyph.middle.d5	0.7872	0.6500	-0.1372
handwritten-neutral-5.encoding.homoglyph.trailing.d1	0.7872	0.5308	-0.2564
handwritten-neutral-5.encoding.homoglyph.trailing.d5	0.7872	0.4990	-0.2882
handwritten-neutral-5.encoding.homoglyph_greek.trailing.d3	0.7872	0.9451	+0.1579

+ V2 evidence validity and unscored observations

Finding groups by validity tier: SILVER 13 · GOLD 1. Mixed-tier groups contribute to each applicable tier; these are group counts.

Review queue · 42 cases

Unresolved REVIEW cases are excluded from automatic vulnerability scoring.

dataset-anger-0.encoding.bidi_reorder.middle; dataset-anger-1.encoding.bidi_reorder.middle; dataset-anger-2.encoding.bidi_reorder.middle; dataset-anger-3.encoding.bidi_reorder.middle; dataset-anger-4.encoding.bidi_reorder.middle; dataset-anger-5.encoding.bidi_reorder.middle; dataset-fear-0.encoding.bidi_reorder.middle; dataset-fear-1.encoding.bidi_reorder.middle; dataset-fear-2.encoding.bidi_reorder.middle; dataset-fear-3.encoding.bidi_reorder.middle; dataset-fear-4.encoding.bidi_reorder.middle; dataset-fear-5.encoding.bidi_reorder.middle; dataset-joy-0.encoding.bidi_reorder.middle; dataset-joy-1.encoding.bidi_reorder.middle; dataset-joy-2.encoding.bidi_reorder.middle; dataset-joy-3.encoding.bidi_reorder.middle; dataset-joy-4.encoding.bidi_reorder.middle; dataset-joy-5.encoding.bidi_reorder.middle; dataset-sadness-0.encoding.bidi_reorder.middle; dataset-sadness-1.encoding.bidi_reorder.middle; dataset-sadness-2.encoding.bidi_reorder.middle; dataset-sadness-3.encoding.bidi_reorder.middle; dataset-sadness-4.encoding.bidi_reorder.middle; dataset-sadness-5.encoding.bidi_reorder.middle; dataset-surprise-0.encoding.bidi_reorder.middle; dataset-surprise-1.encoding.bidi_reorder.middle; dataset-surprise-2.encoding.bidi_reorder.middle; dataset-surprise-3.encoding.bidi_reorder.middle; dataset-surprise-4.encoding.bidi_reorder.middle; dataset-surprise-5.encoding.bidi_reorder.middle; handwritten-disgust-0.encoding.bidi_reorder.middle; handwritten-disgust-1.encoding.bidi_reorder.middle; handwritten-disgust-2.encoding.bidi_reorder.middle; handwritten-disgust-3.encoding.bidi_reorder.middle; handwritten-disgust-4.encoding.bidi_reorder.middle; handwritten-disgust-5.encoding.bidi_reorder.middle; handwritten-neutral-0.encoding.bidi_reorder.middle; handwritten-neutral-1.encoding.bidi_reorder.middle; handwritten-neutral-2.encoding.bidi_reorder.middle; handwritten-neutral-3.encoding.bidi_reorder.middle; handwritten-neutral-4.encoding.bidi_reorder.middle; handwritten-neutral-5.encoding.bidi_reorder.middle;

Diagnostic observations · 126 cases

Diagnostic oracles are unscored regardless of their recognized validity tier. Tier differences are preserved below.

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-anger-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-anger-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-anger-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-anger-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-fear-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-fear-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-fear-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-2.whitespace.repeated_punctuation	REVIEW	Yes

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-joy-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-joy-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-joy-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-joy-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-4.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-sadness-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-sadness-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-sadness-5.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-0.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-0.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-0.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-1.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-1.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-1.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-2.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-2.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-2.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-3.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-3.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-3.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-4.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-4.truncation.signal_head_only	DIAGNOSTIC	No

ATTACK ID	TIER	TIER DIFFERS FROM DIAGNOSTIC
dataset-surprise-4.truncation.signal_tail_only	DIAGNOSTIC	No
dataset-surprise-5.whitespace.repeated_punctuation	REVIEW	Yes
dataset-surprise-5.truncation.signal_head_only	DIAGNOSTIC	No
dataset-surprise-5.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-0.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-0.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-0.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-1.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-1.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-1.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-2.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-2.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-2.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-3.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-3.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-3.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-4.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-4.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-4.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-disgust-5.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-disgust-5.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-disgust-5.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-0.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-0.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-0.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-1.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-1.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-1.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-2.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-2.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-2.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-3.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-3.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-3.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-4.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-4.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-4.truncation.signal_tail_only	DIAGNOSTIC	No
handwritten-neutral-5.whitespace.repeated_punctuation	REVIEW	Yes
handwritten-neutral-5.truncation.signal_head_only	DIAGNOSTIC	No
handwritten-neutral-5.truncation.signal_tail_only	DIAGNOSTIC	No

Evidence bundle

This report is a view of the preserved [analysis.json](#). The [export metadata](#) records SHA-256 hashes of the source and generated files. Keep the original benchmark results, manifests and run metadata with your project evidence.

Source SHA-256

12d47b35c700c6c172ceff5fc071f78952aef7a8695473dc457e9d69229737da

Analysis schema 2 · Report format 1 · Team Checkmate